

Structured Tactile Representations for Bimanual Contact-Rich Manipulation:

Task Topologies, Physical Grounding, and Evidence

XU XIANGYU

PINE Lab / School of Electrical and Electronic Engineering

Nanyang Technological University

Curated scoping-review draft — corpus updated 30 August 2026;
core full-text coding completed 13 August 2026

Abstract

Tactile learning has progressed from task-specific force features to self-supervised encoders, cross-sensor foundation models, spatial anchoring, contact fields, and graph-based dynamics. In parallel, bimanual manipulation has advanced through large-scale simulation, imitation, human-motion retargeting, and increasingly capable real systems. The intersection remains conceptually fragmented: a representation that recognizes local contact is not necessarily one that explains how load, slip, and constraint changes propagate between two hands through an object, tool, workpiece, or environment. This paper presents a curated scoping review of 91 archived records spanning tactile representation, bimanual coordination and handover, relational and physics-informed learning, tool–workpiece control, and datasets. Ninety scholarly works enter the narrative synthesis; publication status is kept explicit because the rapidly moving 2025–2026 frontier contains many preprints and programme-listed papers. The archive-wide evidence map is complemented by full-text precision coding of a purposive 30-paper core (19 Level A and 11 Level B records), with located evidence separated from author claims and reviewer inference.

We organize the literature along six independent axes: task/contact topology, representation scale, structural prior, supervision source, deployment observability, and strength of evaluation evidence. This organization exposes three recurring gaps. First, most tactile encoders model local or temporal content without testing whether cross-interface load paths are recoverable. Second, graph and contact-structured methods often obtain their topology from kinematics, simulator state, or predefined edges rather than infer it from deployable observations. Third, bimanual studies commonly report task success without matched-budget representation baselines, transition-level diagnostics, or topology-level out-of-distribution tests. None of the 30 core papers tested strict held-out contact-topology transfer under a deployment-observable inference protocol. Handover provides a controlled load-transfer laboratory, while cooperative tool–workpiece manipulation adds a harder multi-hop topology. Tool use itself is established, but matched studies of observation-grounded cross-hand load flow remain comparatively sparse in the mapped corpus. We conclude with an evidence-based research agenda centred on observation-grounded structure, interface-level prediction, uncertainty, causal interventions, and transfer across contact topologies. The review does not argue that graphs are intrinsically superior; it specifies the experiments required to establish when explicit structure is useful.

Keywords: tactile representation learning; bimanual manipulation; contact-rich robotics; dynamic contact graphs; handover; load transfer; tool use; visuo-tactile learning; scoping review.

Scope statement. This is a reproducible, curated scoping review of the local research archive, not a PRISMA-complete systematic review or a quantitative meta-analysis. Counts describe the archived corpus and should not be interpreted as publication-rate estimates for the entire field.

1 Introduction

Contact-rich manipulation is partially observable in a particularly consequential way. Vision can estimate object pose before contact, but it cannot directly reveal the pressure redistribution that precedes slip, the direction of an incipient constraint, or whether an apparently stable object is actually supported by the left hand, the right hand, a tool, or the environment. Tactile sensing reduces this ambiguity at an individual interface. Bimanual manipulation then creates a second problem: the controller must relate several local interfaces whose states are coupled through shared rigid-body motion, compliance, friction, and changes in contact topology.

The distinction is easiest to see in a handover. Before transfer, the giver bears most of the load. During dual contact, the same object couples both hands, and a safe release depends on the receiver’s support rather than on giver touch alone. After release, the cross-hand path disappears. A representation that concatenates two tactile streams can, in principle,

approximate this mapping, but it does not state which interfaces interact, which entity mediates the interaction, or why the mapping should transfer to a new grasp topology. Conversely, a graph can name the interfaces, yet a graph whose edges are fixed or supplied by a simulator does not solve contact inference at deployment.

Recent work makes this review timely. Large tactile pretraining systems learn sensor-robust local features from video-based touch, robot play, and multimodal alignment [12, 15, 19, 66, 73]. Spatially anchored and cross-modal models seek to place touch in a geometric or temporal context [17, 22]. Relational models use graphs for tactile arrays, learned dynamics, morphology, and contact physics [1, 3, 13, 48, 68]. At the same time, bimanual policy learning has expanded from simulated suites to human demonstrations and real systems [5, 14, 24, 29, 50]. The 2026 frontier further includes tactile-reactive bimanual control, whole-hand benchmarks, morphology-aware tactile policies, contact-anchored control, and tool-oriented contact fields [11, 23, 36, 42, 70].

The two late-August surveillance updates sharpen this frontier. Hierarchical tactile world models now separate contact, deformation, and slip prediction; sequential visuo-tactile learning uses temporal alignment and training-time privileged posteriors; and physics-aware sensing explicitly models object properties or time-dependent viscoelastic relaxation [34, 56, 61, 64]. Meanwhile, structured cloth dynamics combine a mesh prior with a learned global residual [57]. These advances strengthen the relevant alternatives to a contact graph: a valid comparison is no longer “graph versus concatenation”, but observation-matched structured priors, predictive temporal models, and residual global models. None of these additions, however, directly demonstrates observation-grounded cross-hand load topology.

These advances do not yet answer a shared representation question:

When does an explicit, observation-grounded model of contact topology and load propagation improve bimanual perception or control over an equally capable unstructured temporal model?

This formulation avoids two shortcuts. The first is to equate tactile input with physical grounding. A policy may receive fingertip forces while depending on object state, simulator contacts, or a future reference trajectory. The second is to equate a graph architecture with a scientific contribution. Generic relational learning is well established [3, 55, 69]; the contribution must lie in what the nodes and edges mean, how they can be inferred, which physical quantities they predict, and which counterfactual tests they survive.

1.1 Review objectives and contributions

This review integrates literatures that are often treated separately: local tactile representation learning, structured contact and dynamics models, bimanual manipulation and handover, cooperative tool use, and benchmark design. Its contributions are fourfold.

1. It introduces an orthogonal taxonomy that separates task topology, representation level, structural prior, supervision source, deployment observability, and evaluation evidence. This prevents a method from appearing physically grounded solely because it uses tactile inputs or graph layers.
2. It distinguishes *local contact representation* from *cross-interface coordination representation*. The former encodes what a sensor sees; the latter must preserve which interfaces are coupled, how support is redistributed, and how this relation changes over time.
3. It positions quasi-static handover and cooperative tool–workpiece tasks as complementary experimental topologies. Handover isolates support transfer; tool use tests multi-hop, role-asymmetric coupling without making an unsupported claim that the area is empty.
4. It converts broad calls for “better physical representations” into a falsifiable evaluation programme: matched observation and parameter budgets, oracle, inferred, and random structure comparisons; contact-transition metrics; object-, session-, and contact-sequence-level splits; uncertainty calibration; and targeted interventions.

The intended outcome is not a leaderboard. Sensor form factors, embodiments, training data, task semantics, and success definitions vary too much for a valid pooled effect estimate. Instead, the review identifies which claims can be made from current evidence and which experiments are still needed.

1.2 Terminology

We use *tactile observation* for a measured taxel array, tactile image, deformation field, event stream, or calibrated interface force/torque signal. Simple simulator contact flags and full object state are called *privileged contact variables* unless the deployment system measures them. *Contact topology* denotes the set of active interfaces between hands, objects, tools, workpieces, and the environment. *Load path* denotes a sequence of such interfaces through which forces and moments are transmitted. *Structured representation* covers graphs, contact fields, object-centric frames, spatial anchors, and other models that state relations beyond token order. These terms describe information and assumptions, not a preferred neural architecture.

Table 1: Review questions and the evidence required to answer them.

RQ	Question	Primary evidence fields
RQ1	What must be preserved for bimanual contact-rich manipulation?	Task topology, predicted quantities, temporal horizon, interface identity.
RQ2	How is touch organized across sensors, bodies, entities, and time?	Representation scale, tokenization, spatial frame, structural prior.
RQ3	Is the structure available at deployment?	Sensor inputs, privileged state, inferred versus supplied topology, future reference.
RQ4	How is usefulness demonstrated?	Baselines, split unit, probes, control metrics, OOD tests, interventions, uncertainty.
RQ5	Where are the actionable gaps?	Missing topology–method cells, weak evidence links, unavailable labels or benchmarks.

2 Review Method and Corpus

2.1 Review design

The study is a curated scoping review. Scoping review methodology is appropriate when concepts, evidence types, and evaluation practices are heterogeneous and the goal is to map a field rather than estimate one pooled treatment effect [53]. The accompanying `review_protocol.md` records the scope, eligibility criteria, evidence fields, update procedure, and limitations. The machine-readable `evidence_matrix.csv` is generated from the workspace manifest by a versioned script; rule-assisted annotations remain subject to manual audit. A second evidence layer, stored in `precision_annotations/`, contains full-text precision coding and an individual evidence card for each of 30 purposively selected core papers.

The initial corpus snapshot was frozen on 12 August 2026. It began from a decision-oriented archive assembled for research planning and was expanded around five intersecting areas: (i) tactile and visuo-tactile representation learning; (ii) bimanual dexterous manipulation and handover; (iii) graph, contact-field, object-centric, and physics-informed models; (iv) tool–workpiece contact control; and (v) datasets and benchmarks. Backward and forward chaining was used selectively for landmark and directly competing works. Author project pages, arXiv records, official conference programmes, proceedings pages, and publisher records were used to disambiguate status.

Two rapid-surveillance passes covering 20–26 August and 27–30 August 2026 identified 17 additional eligible studies. These records were screened into the archive-wide layer, their public full texts were retrieved where available, and their implications were integrated into the narrative synthesis. The manuscript corpus was therefore updated and refrozen on 30 August 2026 at 91 records. The 30-paper precision subset was not silently enlarged: the new records have the lighter archive-wide coding described below, and all claims based on the precision subset retain a denominator of 30.

The current draft follows the reporting intent of PRISMA-ScR but does not claim formal PRISMA compliance. It does not reproduce database-by-database search strings, use dual independent screening, or apply a validated cross-design risk-of-bias tool. These omissions matter: the numerical corpus summary is an audit of this archive, not a bibliometric description of the whole field.

2.2 Two-layer evidence design

The archive-wide layer covers all 91 records. Title, abstract, publication and artifact metadata, and selectively consulted full text support eligibility, status accounting, and coarse six-axis coding. This layer provides breadth but is not used as if every row were an independently extracted full-text study.

The precision layer was completed on 13 August 2026 for 30 papers selected for their decision relevance rather than for statistical representativeness. The subset comprises 19 Level A records, coded through full text plus available artifact or code evidence, and 11 Level B records, coded through full-text method and evidence analysis. It covers direct system competitors, executable benchmarks, graph and relational models, contact fields and anchors, mechanics and control references, tool-use systems, and candidate local tactile encoders. Each record distinguishes (i) the authors’ claim, (ii) evidence located in the paper or artifact, and (iii) reviewer inference. It also separates deployment inputs from training-only privilege and codes object, sensor, embodiment, and contact-topology OOD as different claims. All 30 cards contain evidence locations; the procedure is full-text evidence coding, not independent reproduction.




Figure 1: Two-layer corpus flow and publication-status accounting. All 91 records receive archive-wide screening and coarse coding; the purposive 30-paper subset receives full-text precision coding. Counts describe the curated archive, not the global literature.

2.3 Eligibility and coding

Scholarly work was included when it materially addressed robot touch, contact-state or dynamics estimation, tactile-informed control, bimanual coordination, handover, tool use, or an enabling benchmark. Generic graph and physics-learning papers were included only where their formulation directly informs relational contact modelling. A community tactile-data index was retained as a resource record but excluded from study-level synthesis. Duplicate versions were represented by the most informative accessible record, with publication status documented rather than silently upgraded.

Each record was coded along six axes:

1. **Task/contact topology:** single interface, multi-finger hand, two hands sharing an object, sequential/overlapping handover, hand-tool- environment, or hand-tool-workpiece-hand.
2. **Representation scale:** local patch, finger/link, hand, object/interface, cross-entity relation, or trajectory-level latent.
3. **Structural prior:** token order, kinematics, object-centric coordinates, spatial anchors, contact fields/modes, or graph connectivity.
4. **Supervision source:** instance discrimination, reconstruction, cross-modal alignment, action-conditioned prediction, simulator labels, force/pose measurement, or task reward.
5. **Deployment observability:** directly sensed, perception-derived, training-only privilege, or unavailable future reference.
6. **Evidence:** representation probe, open-loop prediction, closed-loop control, real transfer, OOD topology, intervention, and uncertainty.

The axes are kept separate because they answer different validity questions. A real-robot paper can use weak splits; a simulation paper can isolate a causal mechanism cleanly; a high-capacity pretrained encoder can be deployable but offer no cross-interface structure; and a physically interpretable graph can depend on unavailable state.

2.4 Corpus accounting and status discipline

The archive contains 91 records: 14 core or directly competing systems, 24 bimanual and handover works, 19 tactile-representation works, 19 graph, physics, and contact works, eight tool-workpiece works, and seven dataset or benchmark records. Eighty-six PDFs are locally validated; five records retain authoritative source links without a local PDF. Ninety scholarly studies enter the narrative synthesis and one community index is resource-only.

After correcting several manifest entries for identifiable formal proceedings, the evidence matrix records 49 peer-reviewed works, nine works listed as accepted or on an official programme, 32 preprints or author-claim records, and one community resource. “Accepted” is not treated as equivalent to a final publisher record. In particular, rapidly emerging 2026 works are discussed as directions and evidence, not as settled consensus.

2.5 Synthesis strategy

The synthesis is concept-centric rather than paper-by-paper. Representative methods are compared within the same information problem, and negative space is made explicit: for example, a work may demonstrate cross-sensor transfer without testing cross-hand load inference, or closed-loop bimanual success without an explicit tactile representation ablation. No pooled accuracy or success rate is reported because task definitions, sensor density, action spaces, data scale, embodiments, and test splits are not commensurate. Descriptive counts from the precision subset are reported only when the denominator




Figure 2: Two complementary topology families. Handover changes which hand supports a shared object; tool-workpiece manipulation introduces a longer, role-asymmetric load path. Dashed edges are inactive or optional.

and coding rule are explicit. In particular, zero of the 30 core papers evaluated strict held-out contact-topology OOD under deployment-observable inference. Here “strict” means that the participant/edge structure itself is held out; unseen objects, contact modes within a fixed structure, tool geometry, or semantic task compositions do not qualify. Claims that a specific intersection is comparatively sparse are stated as inferences from the mapped corpus rather than field-wide publication counts.

3 Task Topologies and Representation Requirements

3.1 Why topology is a more useful organizing variable than task name

Task labels conceal important physical differences. “Handover” can mean a quasi-static transfer with simultaneous support, a dynamic throw-catch event, or a robot-human release triggered by pull and slip. “Tool use” can mean one hand moving a rigid tool against a fixed environment, one hand stabilizing a workpiece while the other acts, or both hands jointly holding the same tool. These variants demand different information even when their semantic label is the same. A representation-oriented review should therefore begin with the active interfaces and their changes.

Let the task at time t contain entities $\mathcal{V}_t = \{H_L, H_R, O_1, \dots, T_1, \dots, W, E\}$ for left and right hands, objects, tools, a workpiece, and environment. The physical contact topology is $\mathcal{G}_t^c = (\mathcal{V}_t, \mathcal{E}_t^c)$, where an edge represents an active load-bearing or constraint-producing interface. This notation does not require the model itself to be a graph neural network. It provides a neutral way to ask whether the observation history $o_{1:t}$ contains enough information to infer the relevant interfaces and their direction-dependent effects.

3.2 Six topology families

Local single-interface contact. The simplest case associates a tactile patch with local geometry, force, slip, or material. This is the dominant pretraining unit for image-based tactile encoders. It supports transferable perception but leaves global interface identity and cross-body coupling to downstream models.

Multi-finger contact within one hand. In-hand manipulation requires several contacts to be fused across finger kinematics and time. Distributed-tactel graphs and hierarchical tactile graphs explicitly exploit this structure [13, 68]. The object couples the fingers, making this topology a useful precursor to the two-hand problem, but it does not include inter-hand role changes or object transfer.

Two hands sharing one object. Both hands may grasp, reorient, stretch, or stabilize the same object. Bimanual simulation suites and real policy-learning systems cover many such tasks [5, 24, 50]. Success can arise from kinematic coordination and demonstration priors even when load sharing is not explicitly represented.

Handover. Handover is a topology transition rather than merely a motion trajectory. Static and quasi-static transfer offers giver-only, dual-contact, and receiver-only segments; dynamic throw-catch adds a flight phase and impact. The latter is an impressive control test [21] but a weaker primary benchmark for continuous cross-hand load propagation because the hands are not simultaneously connected during flight. Controlled overlapping handover is better suited to studying support redistribution, release readiness, and transition uncertainty [9, 58].

Table 2: Topology-specific representation questions. “Direct” means that an interface is measured locally; “remote” denotes an interface whose effects must be inferred through another entity.

Topology	Critical transition	Observable evidence	Core representation test
Multi-finger in-hand	stick–slip / regrasp	several direct finger contacts	Can local evidence be aligned through kinematics and object motion?
Shared-object bimanual	unilateral–bilateral support	two hand streams, proprioception, optional vision	Is cross-hand dependence captured beyond concatenation?
Quasi-static handover	giver–dual–receiver	load redistribution and contact onsets	Can phase, load share, and safe release be predicted without privileged contacts?
Dynamic handover	release–flight–impact	pre-release contact and post-impact impulse	Does the representation anticipate impact and recovery under temporal discontinuity?
Single-hand tool use	free–working contact	grasp touch plus remote tool response	Can extrinsic tool contact be localized and its mode inferred?
Cooperative tool use	support/actor coupling changes	direct hand contacts; remote tool–workpiece effects	Does structure transfer across a multi-hop, role-asymmetric topology?

One hand, tool, and environment. Tool contact creates extrinsic interfaces whose geometry and compliance may be distant from the tactile sensor. Dynamics models, contact-mode controllers, and contact fields address this setting [18, 44, 51, 54]. It tests whether local hand sensing can identify remote tool contact, but not necessarily whether two hands coordinate.

Two hands with a tool and workpiece. Here one hand actuates a tool and the other stabilizes or repositions the workpiece. Brushing a plate, wiping an object, scrubbing, cutting, peeling, or inserting against a held fixture produces a multi-hop path such as $H_L \rightarrow T \rightarrow W \rightarrow H_R$. BUDS establishes the importance of stabilizer/actor role assignment [14]; TACO documents broad bimanual tool–action–object semantics [35]; and TAMEn includes tactile-aware data collection in tasks such as dish washing [60]. Tool use is therefore an established area rather than an uncrowded task category. The narrower mapped cell—deployment-observable cross-hand load propagation with a held workpiece and matched representation alternatives—contains comparatively few direct studies.

3.3 What a useful representation must retain

Across topologies, five information requirements recur.

1. **Interface identity and quality.** A representation must preserve which finger/link touches which entity, contact confidence, sensor quality, saturation, and missing-data masks. Treating a damaged taxel as zero pressure confounds absence of contact with absence of measurement.
2. **Common spatial meaning.** Local tactile features require a link, hand, object, or task frame before contacts on separate bodies can be compared. Spatial anchoring [22] and object-centric contact fields [18, 36] are different solutions to this frame-assignment problem.
3. **Temporal event structure.** Contact onset, load acceptance, incipient slip, release, impact, and recontact occur on different time scales. A pooled window may recognize a state while obscuring its transition.
4. **Cross-interface dependence.** The representation should allow a downstream model to distinguish coincident local events from a mechanically mediated relation. For handover, receiver loading should predict giver unloading and release readiness; for tool use, tool–workpiece changes should affect the stabilizing hand.
5. **Uncertainty and observability.** Contact edges and remote loads are latent. A credible model should express uncertainty when sensors are occluded, saturated, delayed, or out of distribution instead of returning an overconfident topology.

This topology view also clarifies what is *not* required. A model need not reconstruct every Newton of the complete scene wrench to support safe control; relative load share, future interface wrench, slip risk, or a calibrated contact phase

may be sufficient. Conversely, high task success alone does not show that the learned latent encodes any of these quantities. Representation claims need diagnostic evidence.

4 Tactile Observations, Sensors, and Data Regimes

4.1 The sensor is part of the representation problem

Tactile learning spans RGB or marker-based deformation images, elastomer flow, taxel pressure arrays, event-based tactile cameras, distributed fingertip force, whole-hand skins, and wrist or interface wrench. These signals differ not only in dimension but in invariances, latency, hysteresis, saturation, calibration, and spatial support. A “tactile token” is therefore not a sensor-neutral primitive. Cross-sensor learning must decide which information is shared—contact geometry, normal/shear deformation, dynamics, semantic material cues—and which remains sensor- or embodiment-specific.

Hardware lifetime is part of this invariance problem. A replaceable-cartridge vision-based fingertip reports markedly improved resistance to abrasion and concentrated loading under its accelerated tests [10], while a uniformly illuminated visuo-tactile sensor couples depth reconstruction with spatiotemporal contact-state and slip-direction classification [37]. Durability and illumination stability do not by themselves solve representation transfer, but degradation state, cartridge identity, and lighting/calibration changes should be recorded as sensor-domain variables rather than absorbed into unexplained model error.

Vision-based tactile sensing has enabled large-scale self-supervision because a frame can be processed with mature visual backbones. Sparsh learns general touch representations from large unlabeled tactile collections [19]; T-Dex uses robotic play to pretrain features useful for dexterous control [15]; UniTouch and AnyTouch align heterogeneous sensors and static/dynamic content [12, 66]; and T3 studies transferable tactile Transformers [73]. These works make local tactile content more reusable, but most cross-sensor objectives do not preserve a global hand-object-hand contact identity by themselves.

Distributed and whole-hand sensors shift the difficulty. A fingertip array has a natural patch boundary; a skin spanning phalanges and palm requires morphology and sensor-layout metadata. HTT explicitly addresses heterogeneous tactile sensors [4]; HT-Bench expands full-hand tactile representation with egocentric vision [23]; FTP-1 introduces morphology-aware tactile tokens across sensors and embodiments [70]. Such work is directly relevant to bimanual transfer because the two hands may have different coverage, calibration, or missing channels.

4.2 Force estimates, measured wrench, and ground truth

Three quantities are frequently conflated:

- a raw tactile deformation or taxel field;
- a wrench estimated from that tactile signal by calibration or a learned model; and
- an independently measured interface wrench from a dedicated force/torque sensor or instrumented object.

The distinction determines what can serve as a target. Training a model to predict an estimate derived from the same tactile image measures compression or consistency, not independent physical validity. Likewise, a simulator’s net fingertip contact force is useful supervision but is not equivalent to a real tactile array. PhysGraph, for example, studies graph-structured bimanual policy learning with simulator state and fingertip resultant forces rather than raw tactile images [30]. This makes it relevant to relational control while limiting conclusions about tactile representation transfer.

Force inference can also drift even when the contact geometry appears unchanged. Relaxation-aware soft-gripper sensing uses onboard vision, infrared thermography, and a temperature-coupled viscoelastic model to estimate force during sustained holds [56]. It shows why contact duration, temperature, and loading history belong in the observation audit. An instantaneous tactile or force frame may be insufficient for a long hold, and agreement with a same-frame proxy should not be mistaken for relaxation-robust force estimation.

For load-transfer research, the strongest target hierarchy is:

1. independent instrumented-object or interface force/torque where feasible;
2. calibrated robot wrist/hand wrench with explicit coordinate transforms and gravity compensation;
3. simulator wrench for pretraining or privileged-to-observation distillation, followed by a real validation subset;
4. proxy load share or contact phase, with the proxy’s construction and circularity documented.

Table 3: Representative data regimes and what they can and cannot establish for cross-interface representation. Status reflects the corpus freeze.

Resource / line	Regime	Main strength	Caution for bimanual load-path claims
Touch and Go [65]	human-collected vision–touch	broad multimodal diversity	limited robot action, proprioception, and calibrated cross-hand load labels
T-Dex / Sparsh [15, 19]	robot touch / large unlabeled tactile data	reusable local encoders	downstream topology not specified by pretraining objective
VTDexManip [33]	simulated visual–tactile benchmark	reproducible pretraining and handover task	sim observation and contact fidelity must match intended deployment
T-Rex [42]	large real robot tactile data	temporally rich bimanual control evidence	dataset subset, synchronization, and exact labels require audit
OAKINK2 / TACO [35, 71]	bimanual human/object datasets	topology and semantic task breadth	tactile load evidence is not their primary signal
HT-Bench / TactiDex [23, 40]	whole-hand real tactile benchmarks	coverage beyond fingertips	recent status and resource release should be verified per artifact
ViTacPhys [34]	human visual–tactile demonstrations and robot transfer	explicit mass, friction, and stiffness probes	semantic priors and unreleased artifacts must be separated from tactile evidence
SoftVTBench [25]	deformable-object visual–tactile benchmark	independent deformation state and safety-aware success	public release version must be matched to the reported manuscript scale
Tactile Genesis [8]	scalable tactile simulation	layout variation and dense labels	sim-to-real contact appearance and dynamics gap

4.3 Data regimes and their affordances

Human-collected datasets can provide semantic and visual diversity at scale. Touch and Go aligns human-collected touch and vision [65]; OAKINK2 captures complex bimanual hands–object activity [71]; and TACO focuses bimanual tool–action–object understanding [35]. Their strength is breadth, while precise synchronized robot proprioception, calibrated loads, and intervention labels may be absent.

Robot-collected data provides action, embodiment, and timing. VTDexManip offers a visual–tactile pretraining and dexterous manipulation benchmark that includes bimanual handover [33]. T-Rex reports a large tactile-reactive robot dataset and bimanual capabilities [42]. RCT targets robot-collected touch–vision–language generalization [16], while TactiDex and HT-Bench expand real, whole-hand tactile evaluation [23, 40]. As of the artifact audit on 8 August 2026, the TactiDex author project page stated ACM MM 2026 acceptance, but no downloadable dataset, code, or model artifact had been confirmed; it should not be assumed to be an executable first-month substrate. Because several of these appeared as preprints or accepted works near the corpus freeze, exact public availability, license, and final metadata must be checked before they are treated as reproduction substrates.

Two new datasets illustrate complementary physical labels. ViTacPhys uses human visual–tactile demonstrations to estimate mass class, friction class, and stiffness before conditioning an adaptive grasping policy [34]. SoftVTBench synchronizes policy-visible visual–tactile signals with evaluator-only finite-element state and defines a deformation-aware success criterion [25]. The former makes physical-property probes concrete; the latter shows that task completion can co-exist with excessive deformation. Artifact status remains part of the evidence: at the 30 August audit, ViTacPhys code/data were still listed as forthcoming, and the public SoftVTBench release described a smaller dataset version than the current manuscript.

Simulation is essential for contact labels and topology interventions. Tactile Genesis targets scalable tactile simulation across layouts [8]; generic graph simulators demonstrate how object/particle relations can support long-horizon physical prediction [2, 48]. Simulation can generate edge, force, slip, and counterfactual labels that are prohibitively expensive in the real world. Its risk is an unobserved shortcut: a model may use precise object state, collision geometry, contact impulses, or future reference that will not exist at deployment.



Figure 3: Representation hierarchy. The arrows denote additional organization, not a mandatory neural architecture. Cross-interface claims require evidence at the rightmost levels, even when the local encoder is strong.

4.4 Synchronization is not a preprocessing detail

Bimanual tactile data is intrinsically multi-rate. Cameras, tactile arrays, joint encoders, wrist wrenches, and policy actions may run at different rates and have different transport delays. A 50–100 ms misalignment can change an apparent causal relation at contact onset or release. Dataset construction should therefore retain raw timestamps, clock source, interpolation policy, sensor-quality masks, and commanded versus measured action times. Synchronization should be evaluated with event-based sanity checks—for example, whether tactile onset precedes an object acceleration that it could physically cause—rather than only by visually aligned plots.

The split unit is equally important. Frame-level random splits leak object, episode, grasp, and contact-sequence identity into the test set. Representation claims require episode-level splits at minimum, with object, session, operator, sensor, embodiment, and contact-sequence holds-outs selected according to the claim. For cross-topology transfer, the complete topology family—not just a motion instance—must be absent from training.

4.5 A minimum interoperable schema

A useful cross-dataset schema should keep (i) raw tactile observations and quality masks; (ii) calibrated per-sensor meta-data; (iii) joint states and link poses; (iv) object/tool/workpiece identifiers and optional perception-derived poses; (v) actions and timestamps; (vi) phase, interface, slip, and wrench labels with provenance; (vii) temperature, contact duration, maintenance/calibration state, and material metadata when relaxation or wear matters; and (viii) split membership. Coordinate frames must be explicit. Without this information, a study can compare end-to-end policies, but it cannot reliably compare structured interface representations.

5 From Local Touch to Cross-Interface Representations

5.1 A hierarchy of representation scales

Tactile models can be placed on a hierarchy from sensor-local appearance to task-level physical relations ([figure 3](#)). Progress at a lower level is enabling but does not imply progress at every higher level. For example, a sensor-invariant embedding may improve local contact classification while discarding absolute orientation or force magnitude that a load-transfer model needs. Conversely, a task-specific relational policy may model useful interactions while inheriting a fragile local encoder.

5.2 Local self-supervised and transferable touch

Self-supervised tactile learning replaces expensive per-task labels with objectives such as reconstruction, masked prediction, temporal consistency, and instance discrimination. T-Dex couples pretraining to robotic play, showing that unlabeled interaction can benefit downstream dexterity [15]. Sparsh scales self-supervision over vision-based tactile data and evaluates a family of downstream tasks [19]. These approaches establish that generic visual backbones are not enough and that tactile-specific pretraining can be reused.

Transfer across sensors is harder because pixels encode different elastomers, illumination, marker patterns, optics, and spatial scales. UniTouch learns unified multimodal tactile representations [66]; AnyTouch explicitly addresses static and dynamic information across multiple visuo-tactile sensors [12]; T3 studies transferable Transformers [73]; and UniT uses a unified tactile modelling approach [62]. Sparsh-X broadens the multisensory alignment setting [20], while TactAlign maps human tactile-glove features into a robot tactile space for cross-embodiment transfer [59]. The common scientific question is not simply whether embeddings transfer, but which tactile factors remain identifiable after alignment. A sensor-invariant latent that collapses shear direction or scale may be detrimental to load-path reasoning.

Whole-hand and morphology-aware models extend the unit of representation. HTT uses heterogeneous tactile tokenization and masked learning [4]; FTP-1 incorporates sensor morphology in a large tactile policy [70]; HT-Bench supplies a whole-hand benchmark with egocentric vision [23]. These works motivate explicit metadata for sensor location, normal

direction, link identity, and quality. In bimanual systems, morphology-awareness must also prevent a shortcut in which “left” and “right” sensor IDs stand in for task role.

5.3 Temporal tactile representations

Touch is an event stream. Static contact appearance can be ambiguous: the same deformation may correspond to loading or unloading, stable stick or the moment before slip, depending on recent history. TactileVAD highlights geometric aliasing in learned tactile dynamics [43]. Reactive control systems increasingly split slow semantic reasoning from fast tactile reaction; T-Rex uses tactile-reactive control at different time scales [42], while recent slow-fast visuo-tactile diffusion policies pursue a related control principle. ViTacFormer also includes future-touch prediction as part of cross-modal representation learning [17].

HiTac-WAM makes the prediction target more explicit: candidate actions are evaluated through a directed contact $\rightarrow$ deformation $\rightarrow$ slip hierarchy, and persistent forecast error can trigger replanning [64]. VT-MUSE instead learns a frozen sequential visuo-tactile representation using cross-modal temporal alignment, masked-view consistency, and a training-only posterior that sees the full tactile history [61]. These are stronger controls than a single future-frame decoder. They also expose distinct hypotheses: whether physical variables should be ordered explicitly, and whether a deployable prior can retain information learned from richer training-time histories.

Temporal modelling choices include fixed windows, recurrent states, temporal convolution, causal Transformers, and state-space models. Their comparison is often confounded by receptive field, parameter count, and action history. For contact-transition evaluation, causal alignment is essential: a future tactile frame may be a valid training target, but it must not leak into the deployed encoder input. Reporting should distinguish synchronous state estimation, 100–300 ms forecasting, and closed-loop reaction.

Vision/proprioception-only models also strengthen the temporal null hypothesis. TemporalFlow-VLA supervises ordered execution-history queries with training-only robot-surface flow [67]; PredVLA updates a compact recurrent latent through prediction-error inference [49]. Neither supplies distributed touch, so neither closes the tactile relation gap. They are nevertheless important matched baselines: a new tactile structure should outperform comparable history grounding and prediction-error correction, not merely a memoryless visual policy.

5.4 Multimodal fusion and spatial anchoring

Tactile measurements acquire meaning from proprioception and vision. Early or late concatenation is a strong baseline because a sufficiently expressive temporal model can learn cross-modal interactions. Cross-attention gives tactile tokens selective access to image, state, or action tokens; ViTacFormer is a prominent recent example [17]. The fair question is whether a proposed structure improves over this baseline at equal observation, compute, and temporal context.

Physical-property fusion adds another alternative. ViTacPhys predicts mass, friction class, and stiffness from human visual-tactile demonstrations and uses the estimates to condition grasping [34]. Such property tokens can explain adaptation across objects without representing a cross-hand load path. A fair system-level comparison should therefore include both property- conditioned fusion and relation-conditioned fusion, and ask which failure modes each resolves.

Spatial anchoring transforms local touch into a shared frame. SaTA anchors tactile awareness using forward kinematics, enabling spatial reasoning beyond sensor index [22]. Object-centric approaches similarly normalize motion with respect to the manipulated entity [6]. Anchoring is particularly valuable for two hands because corresponding tactile features lie in separate kinematic trees. Yet anchoring assumes that the relevant transform is known or perceived; errors in link calibration and object pose should be tested, not hidden.

5.5 What existing objectives supervise

Pretraining objectives privilege different information. Reconstruction preserves pixel detail but may waste capacity on illumination or marker appearance. Masked prediction encourages context use but can exploit spatial redundancy. Contrastive alignment emphasizes shared content but can suppress modality-specific physical signals. Future prediction encodes dynamics but may average multimodal outcomes. Task reward discovers control-relevant features but can exploit spurious state or demonstration regularities. No objective is universally correct; its sufficiency must be judged against the intended quantities and topology.

For bimanual coordination, promising auxiliary targets include contact onset and release, future interface occupancy, load-share ratio, future local wrench, slip risk, and calibrated uncertainty. These are not merely extra labels. They test whether the latent exposes a mechanism that a downstream policy could use. Where labels are available only in simulation, privileged-to-observation distillation can train an observable student, but evaluation must verify that the student does not require privileged inputs at test time.

Table 4: Representation families, what they organize, and the unresolved cross-interface question. Entries are representative rather than exhaustive.

Family	Representative works	Primary organization	Unresolved bimanual question
Local SSL	T-Dex, Sparsh [15, 19]	contact appearance, geometry, dynamics	Does the latent retain load magnitude/direction and interface identity?
Cross-sensor alignment	UniTouch, AnyTouch, T3, Sparsh-X, TactAlign [12, 20, 59, 66, 73]	shared tactile factors across hardware	Which physically relevant factors are lost to invariance?
Whole-hand / morphology-aware	HTT, HT-Bench, FTP-1 [4, 23, 70]	sensor layout, link/body coverage	Can roles and topology transfer across two heterogeneous hands?
Temporal dynamics	TactileVAD, T-Rex, HiTac-WAM [42, 43, 64]	event history, directed tactile state, and action-conditioned change	Are transitions predicted causally, and does the directed hierarchy beat a matched flat head?
Sequential cross-modal	ViTacFormer, VT-MUSE [17, 61]	vision–touch–state interaction and training/deployment temporal latents	Does explicit interface structure add value to matched sequential fusion?
Physical-property latent	ViTacPhys [34]	mass, friction, and stiffness estimates	Do property tokens explain adaptation without cross-interface relations?
Spatial anchoring	SaTA [22]	tactile features in kinematic space	How robust is the representation to pose error and topology change?

5.6 Necessary baselines for representation claims

A new structured encoder should be compared against more than raw concatenation. At minimum, a matched study should include: proprioception-only; tactile-only; concatenation with the same local encoder; cross-attention; a temporal Transformer or recurrent model; a prediction-error recurrent model; a flat multi-head and a directed tactile hierarchy; training-only temporal grounding; spatial anchoring where poses are available; a static graph; the proposed inferred structure; and an oracle structure used only as an upper bound. Parameter count, input history, optimization budget, augmentations, and pretraining data must be controlled or reported. Without these controls, an improvement may come from capacity, temporal context, objective ordering, or privileged supervision rather than representation structure.

6 Graphs, Contact Fields, and Physical Grounding

6.1 A graph is a hypothesis about factorization

Interaction Networks established a general relational inductive bias: entities and their pairwise relations can be modelled separately, with messages aggregated to update entity states [3]. NerveNet applies graph policies over articulated morphology [55]; graph simulators learn physical dynamics over interacting particles or bodies [48]; Graphormer demonstrates that Transformer attention can be modified with graph-structural encodings [69]; and AnyMorph studies policies across morphologies [52]. Together these works make “uses a GNN” an architectural choice, not a sufficient novelty claim.

For tactile bimanual manipulation, a graph encodes a stronger hypothesis: task dynamics factorize around interfaces and the paths connecting them. This hypothesis is useful only if (i) node and edge semantics correspond to deployable or inferable quantities, (ii) the factorization improves sample efficiency, robustness, interpretability, or intervention response, and (iii) the same gain cannot be explained by additional state or parameters.

6.2 Taxel, finger, and morphology graphs

Graph construction at the tactile-array or finger level is relatively direct. Distributed tactile sensors have known positions on an articulated hand; convolution or message passing can share information along this geometry. Multi-finger in-hand manipulation with distributed tactile GCNs demonstrates this idea [13]. TacGNN uses hierarchical graph structure for blind tactile-based in-hand manipulation [68]. TactiGraph treats event-based tactile measurements asynchronously for

contact-angle prediction [47], and TouchWGNN develops spatio-temporal tactile graphs for multimodal dexterous manipulation [41].

These models address meaningful spatial and temporal organization. However, a taxel adjacency graph is not the same as a dynamic task contact graph. Taxels can remain neighbours in the sensor layout whether or not their links contact the same object; fingers can be kinematically adjacent while transmitting loads through different entities. Bimanual load-path work therefore needs at least two levels: relatively stable morphology/sensor structure and dynamic inter-entity interfaces.

6.3 Dynamics graphs and contact discontinuities

RoboPack learns tactile-informed recurrent graph dynamics for dense packing and uses the model for planning/control [1]. Its relational factorization is compelling because tool, objects, and environment interact over time. ContactNets targets discontinuous contact dynamics with a structured implicit model [46]; graph network simulators have been shown to model rigid contact transitions [2]. This line of work suggests two design principles for bimanual representation: predict changes at interfaces, and avoid treating discontinuity as ordinary smooth motion noise.

MeshPriorDiT adds a useful hierarchical counterexample to an architecture-only comparison. An action-conditioned mesh GNN supplies topology and grasp-constraint priors, while a flow-matching diffusion Transformer models residual and global cloth effects [57]. Its evidence concerns deformable mesh dynamics rather than tactile load transfer, but the factorization is directly relevant: structured local priors and expressive global residuals can be complementary. A bimanual study should therefore compare graph-only, global-residual-only, and combined models rather than forcing a false GNN-versus-Transformer dichotomy.

Yet dynamics prediction can be correct for the wrong reason. Exact object pose, collision geometry, and contact labels can make a simulator graph accurate while an observation-only system remains underdetermined. A useful paper must specify which signals construct the graph at training and at deployment. “Physics grounded” should describe targets, conservation constraints, frames, or interventions—not merely that the task occurs in a physics engine.

6.4 Contact fields, modes, and anchors

Graphs are not the only structured representation. Neural Contact Fields track extrinsic contact through tactile sensing [18]; object-centric contact-field approaches extend this spatial view to visuo-tactile state estimation and tool interaction. Semantic Contact Fields predict semantically meaningful contact/force structure for tool manipulation [36]. Contact-mode control represents discrete constraint regimes for non-prehensile manipulation [44]. Contact-anchored policies and contact-grounded policy learning similarly use contact as a privileged or learned organizing variable [11, 63].

Recent work also clarifies two meanings of “contact structure.” One-shot learning by identifying physical interactions extracts making/breaking contact events from a demonstration and uses them as transition conditions for hybrid position–force control [45]. CoToGrasp conditions grasp synthesis on a requested contact topology in a canonical workspace [39]. Both make structure explicit, but the former uses a demonstrated event state machine and the latter a prescribed grasp-contact layout. Neither is equivalent to inferring a time-varying cross-hand topology from raw deployment observations. They are valuable controls precisely because they test how much structure can be specified rather than learned.

Fields are attractive when contact location varies continuously over object surfaces; discrete graphs are attractive when entity/interface identity and message routing matter. Hybrid representations can use a field to infer contact location and confidence, then instantiate an entity-level relation. The decisive comparison should be empirical: a field, graph, anchored token set, and unstructured attention model should receive the same observations and predictive targets.

6.5 PhysGraph as a boundary case

PhysGraph is a directly relevant early graph-policy formulation for bimanual hand–tool–object manipulation [30]. It organizes hand, tool, and object information in a graph Transformer and reports strong simulated control. Its input regime, however, includes simulator state, fingertip resultant forces, object/tool geometric encodings, and future reference. It should therefore be read as evidence for relational policy factorization under privileged state, not as tactile-array representation pretraining or observation-grounded contact inference.

A local audit of the public main-branch snapshot at the corpus freeze found that the bias generator exposes contact-mask and geometric-position inputs, while the default forward call does not pass them; the resulting public path uses the configured topology rather than the advertised dynamic contact/geometric bias. This is a reproducibility observation about that code snapshot, not evidence that the authors’ reported experiments used the same path. The correct experimental response is to reproduce both the stock public path and a clearly documented corrected-bias path, then avoid attributing their difference to the paper without author confirmation.

This case illustrates why code-level observability matters. A diagram can show dynamic edges while the executed model uses static connectivity; a paper can describe tactile feedback while the input is a low-dimensional simulator force; and a physically plausible target can be constructed from unavailable future state. Reproduction should trace every tensor from sensor or simulator source to the model and record its deployment availability.

Table 5: Structured-representation choices. The final column states what must be demonstrated rather than assumed.

Structure	Natural unit	Typical source	Validity test
Sensor/morphology graph	taxel, link, joint	calibrated layout and kinematics	robustness across layouts, missing sensors, and embodiments
Static task graph	hand, object, tool	designer-specified entity relations	advantage over matched attention without dynamic edge changes
Dynamic contact graph	active interface	tactile/proprioceptive inference or simulator contact	inferred–oracle gap, transition timing, false-edge interventions
Object-centric anchor	local feature in object frame	visual/kinematic pose estimate	degradation under pose error and unseen grasps
Contact field	continuous surface point/force	tactile/vision and geometry	contact localization, force consistency, surface/OOD generalization
Contact mode	discrete constraint regime	estimator, demonstration, or privileged label	mode-transition accuracy and closed-loop recovery
Physics residual/constraint	force balance, momentum, complementarity	measured or simulated quantities	calibration, counterfactual consistency, no label circularity

6.6 Physical objectives for cross-interface structure

Suppose an inferred interface edge $e_{ij,t}$ has probability $p_{ij,t}$ and an estimated wrench $\hat{w}_{ij,t}$ expressed in a common object frame. Useful objectives can combine:

$$\mathcal{L} = \lambda_e \mathcal{L}_{\text{edge}} + \lambda_w \mathcal{L}_{\text{future-wrench}} + \lambda_s \mathcal{L}_{\text{slip}} + \lambda_p \mathcal{L}_{\text{phase}} + \lambda_c \mathcal{L}_{\text{consistency}}. \quad (1)$$

The consistency term might penalize disagreement between predicted interface wrenches and measured object acceleration or enforce action–reaction relations when both sides of an interface are observed. It should be soft and uncertainty-aware: unmodelled compliance, gravity-compensation error, and unobserved environmental contact make exact balance inappropriate. The model should also predict confidence, because a missing edge under sensor occlusion is not equivalent to a confidently absent contact.

For sustained contact, the state should additionally represent slow relaxation. Temperature-coupled viscoelastic force estimation shows that identical-looking instantaneous configurations can correspond to different force histories [56]. Interface objectives should therefore be evaluated at multiple contact durations and, where relevant, condition on temperature or an identified relaxation state rather than treating drift as independent noise.

The central empirical question remains comparative. If a temporal Transformer with the same inputs predicts future wrench and phase equally well, explicit edges may add interpretability but not predictive value. If structure improves data efficiency, transition accuracy, or transfer to a new topology—and the benefit disappears under random or time-shuffled edges—the relational hypothesis receives stronger support.

7 Bimanual Coordination, Handover, and Cooperative Tool Use

7.1 Bimanual learning is broader than tactile coordination

Bimanual manipulation has grown rapidly through simulation benchmarks, retargeting, imitation, and scalable demonstration generation. Bi-DexHands provides a broad simulated reinforcement-learning benchmark [5]. Bimanual Dexterity for Complex Tasks and related real systems demonstrate increasingly capable coordinated policies [50]. DexMimicGen generates bimanual data from demonstrations [24]; ManipTrans transfers behaviours through residual learning [29]; object-centric human-motion approaches organize demonstrations around manipulated objects [6]. These works establish task and data scalability. They do not automatically test whether a tactile latent explains cross-hand forces or contact transitions.

Late-August work further raises the vision/proprioception-only baseline. Training-time modality masking improves robustness of a query-based bimanual VLA without adding a tactile relation model [7]; PartialBiGrasp infers hidden local geometry from partial point clouds and searches for force-closure-compatible grasp pairs [26]. CLAP learns a cross-embodiment action-conditioned video world model with bimanual data among its supported settings [32], while rapid on-robot juggling learns local dynamic models around a global prior under a reachable safe-set constraint [28]. Together they explain more bimanual success through visual geometry, robust fusion, video prediction, and rapid dynamics adaptation. A tactile representation must therefore demonstrate an incremental benefit under occlusion, friction change, slip, or

simultaneous load transfer, not rely on bimanual task success alone.

This distinction matters for literature positioning. The broad claim “bimanual manipulation is underexplored” is no longer defensible. The narrower intersection of *observation-grounded tactile structure*, *explicit interface/load targets*, and *matched representation comparisons* remains much less populated in the mapped corpus. A contribution should therefore be framed around that intersection, not around the existence of two hands.

Bi-Touch is an important early direct baseline: it used a concatenated tactile policy state and sim-to-real reinforcement learning for contact-rich bimanual pushing, reorientation, and gathering tasks [31]. It shifted the question from whether touch can help bimanual control to whether an explicit relational representation adds value beyond a matched tactile policy input. A new graph method should therefore include an equivalent concatenation baseline even if the original Bi-Touch implementation is not the primary reproduction target.

7.2 Roles, symmetry, and mechanical necessity

A bimanual taxonomy can distinguish symmetric and asymmetric actions, shared and separate objects, coordinated and independent motions, and temporal role changes [27]. For representation research, two additional questions are decisive:

1. **Is coordination mechanically necessary?** If one hand can complete the task while the other supplies a decorative or easily ignored motion, task success does not diagnose cross-hand representation.
2. **Can the roles be exchanged?** A representation that encodes fixed left/right identities may fail when actor and stabilizer swap, even though the physical relation is unchanged.

BUDS makes stabilizer–actor coordination explicit and demonstrates that one hand can create conditions for the other [14]. This is a useful conceptual bridge from shared-object manipulation to tool–workpiece tasks. However, role assignment should be coupled to measurable necessity: removing, delaying, or perturbing the stabilizing hand ought to change forces, error, or success in a predictable way.

7.3 Handover as a load-transfer benchmark

Handover provides an unusually clean phase structure. Human-inspired and robot–robot/robot–human controllers have long used load, pull, grip force, or slip cues to govern transfer and release [9]. Contact-aware dexterous handover further emphasizes when and where contacts occur [58]. VTDexManip supplies a reproducible simulated bimanual handover task with visual-tactile observations [33]; recent tactile-reactive systems expand the scale and realism of bimanual touch [42].

The most informative handover interval is dual contact. Let $F_L(t)$ and $F_R(t)$ be gravity-axis support estimates expressed in a common object frame. A simple load-share target is

$$\rho_R(t) = \frac{\max(0, F_R(t))}{\max(0, F_L(t)) + \max(0, F_R(t)) + \epsilon}. \quad (2)$$

This ratio is interpretable and less scale-sensitive than absolute force, but it is only valid if the force estimates are independently calibrated and frame aligned. Contact phase, time-to-support, future receiver wrench, giver slip risk, and release readiness provide complementary targets. The point is not to impose one scalar definition; it is to make the transfer mechanism measurable.

Dynamic throw–catch handover explores timing, ballistic motion, and impact recovery [21]. It should be a later robustness task rather than the primary cross-hand load benchmark, because the flight phase removes the simultaneous object-mediated path. A quasi-static transfer with a speed perturbation, object-mass change, or receiver delay offers cleaner causal tests first.

Robot juggling reinforces this sequencing: impressive rapid online adaptation in a dynamic bimanual task can be achieved from onboard vision and a local dynamics model [28]. It is therefore a strong adaptation reference, but its flight/contact alternation does not provide the persistent, simultaneous object-mediated load path needed to validate cross-hand tactile structure.

7.4 Cooperative tool–workpiece manipulation

One hand holding a brush while the other holds a plate is closely aligned with the same representation problem, but its topology is more demanding. The actor hand senses its grasp on the tool; the stabilizer hand senses the workpiece; the working interface between brush and plate is remote from both local grasp interfaces. Force and motion propagate through at least three contacts. This task therefore tests whether a representation can move from $H_L \leftrightarrow O \leftrightarrow H_R$ to $H_L \leftrightarrow T \leftrightarrow W \leftrightarrow H_R$.

The adjacent literature supplies components rather than a saturated direct benchmark. Tactile Tool Manipulation learns tool-related control with tactile feedback [51]; compliant tool–environment dynamics and extrinsic contact-mode control model the working interface [44, 54]; RoboPack demonstrates tactile-informed relational dynamics in tool-mediated dense

Table 6: How the two principal task families complement each other. “Mapped evidence density” is a qualitative inference from the curated corpus, not a bibliometric claim.

	Quasi-static handover	Cooperative tool–workpiece use
Core topology	hand–object–hand, changing support	hand–tool–workpiece–hand, persistent role asymmetry
Primary signal	load redistribution and release	working–contact wrench, stabilizer reaction, slip
Main advantage	clean phases and interpretable transition	harder multi-hop transfer and practical task relevance
Main confound	object pose/motion can reveal phase	visual trajectory, periodic motion, and compliance can reveal task state
Mapped evidence density	established handover/control literature; few matched representation studies	established tool/contact systems; few matched cross-hand load-flow studies
Best research role	first controlled representation benchmark	gated second topology after mechanism is established

contact [1]; TACO covers tool–action–object semantics [35]; TAMEn targets tactile-aware closed-loop data collection in contact-rich tasks such as dish washing [60]; and Semantic Contact Fields make tool contacts explicit [36]. These works make the area adjacent, not empty. What remains scarce is a controlled study of *cross-hand tactile load propagation* with a held workpiece and matched representation baselines. TAMEn is a direct system competitor for a dish-washing demonstration; novelty cannot rest on performing that task with touch, but must come from observation-grounded load-path inference, mechanism-level evidence, or cross-topology transfer.

7.5 Designing a diagnostic tool task

A diagnostic task should make the second hand necessary while remaining simple enough to label. A rigid stylus or wiping puck against a plate is preferable to a deformable brush for an initial experiment: the contact location, normal, and applied path can be controlled, and simulation is more credible. The task can then progress to a compliant brush once sensing and analysis are stable.

Useful experimental factors include tool stiffness, workpiece mass, friction, path curvature, contact-force target, stabilizer grasp, actor/stabilizer role swap, and brief perturbations of either hand. A *coordination-necessity check* should compare two hands, fixed workpiece, passive second hand, and removed second-hand tactile input. If the actor succeeds with no measurable cost when the stabilizer is removed or frozen, the task is not a strong cross-hand benchmark.

The tool task should enter a research programme after the simpler handover study has established three points: inferred structure improves at least one transition/load target over matched unstructured baselines; the gain is not due to privileged state; and structure remains useful under object/grasp holds-out. Before those conditions, tool use is valuable as a frozen task specification and offline proxy, but a full hardware campaign risks combining unsolved sensing, calibration, control, and representation problems.

7.6 Cross-topology transfer as the stronger claim

Training and testing separately on handover and tool use would produce two task results. A stronger representation claim asks whether a model pretrained to infer interfaces and loads on handover transfers to the longer tool topology with fewer labels or demonstrations. The test should distinguish:

- **encoder transfer:** freeze the local tactile encoder and retrain only the relation head;
- **relation transfer:** preserve edge functions while instantiating a new entity topology;
- **full fine-tuning:** initialize all weights and measure data efficiency;
- **zero-shot probing:** test whether phase/load variables are linearly recoverable without task fine-tuning.

If only full fine-tuning helps, the model may provide generic initialization rather than compositional topology transfer. If relation modules transfer under role swap and added tool/workpiece nodes, the evidence for structured factorization is substantially stronger.




Figure 4: An evidence ladder for representation research. Higher rungs depend on, rather than replace, the lower-level checks.

8 Evaluation: From Task Success to Representation Evidence

8.1 Why task success is insufficient

Closed-loop success is necessary for a control method but weak evidence about a specific latent mechanism. A policy can infer phase from vision, memorize object trajectories, exploit a scripted receiver, or rely on privileged object state. It may succeed without using tactile inputs at all. Conversely, a representation can predict contact transitions accurately while a poorly tuned controller masks the benefit. Evaluation should therefore connect four levels: observation integrity, representation probes, open-loop dynamics, and closed-loop behaviour.

8.2 Matched-budget baseline matrix

The baseline matrix in [table 7](#) isolates the source of an improvement. All learned variants should share the same local tactile encoder, observations, history, training examples, optimizer budget, and approximately matched parameter/compute count. If a graph variant receives link poses while a Transformer does not, the study tests information access, not factorization. Two additional controls are necessary after the core-paper audit. A predictive contact-grounding baseline tests whether future tactile/state objectives alone explain the gain [\[17, 63\]](#); a dynamic taxel graph tests whether local sensor-level message passing already suffices [\[41\]](#). The surveillance update adds three more controls: a directed contact–deformation–slip predictor [\[64\]](#), training-only temporal-flow grounding [\[67\]](#), and prediction-error recurrent inference [\[49\]](#). Cross-embodiment visual prediction such as CLAP is a useful system-level baseline where video pretraining is available [\[32\]](#), but it must receive the same deployment observations when used to isolate the value of touch.

Table 7: Minimum matched-budget experiment matrix. The oracle graph is an upper bound and must never be presented as a deployable method.

Model	Same local encoder	Temporal context	Explicit structure	Question isolated
Proprioception-only	n/a	matched	none	Is touch needed?
Tactile concat	✓	matched	none	Does local touch help?
Cross-attention	✓	matched	implicit	Is flexible token interaction sufficient?
Temporal Transformer	✓	matched	implicit	Is time/context the main source of gain?
Prediction-error recurrent	✓	matched	implicit/predictive	Does observation-driven latent correction explain the gain?
Predictive grounding	✓	matched	implicit/predictive	Do future tactile/state objectives explain the gain?
Directed tactile hierarchy	✓	matched	ordered physical heads	Does contact → deformation → slip ordering beat flat heads?
Training-only temporal grounding	✓	matched	privileged supervision only	Does richer training supervision, rather than deployed structure, explain the gain?
Spatially anchored tokens	✓	matched	coordinates	Does a shared frame suffice?
Dynamic taxel graph	✓	matched	local learned edges	Does sensor-level graph structure suffice?
Static graph	✓	matched	fixed edges	Does relational routing help without edge inference?
Inferred dynamic graph	✓	matched	observed/inferred	Does topology inference add value?
Oracle graph	✓	matched	privileged edges	What is the attainable structural upper bound?
Random / shuffled graph	✓	matched	invalid control	Does edge semantics, rather than regularization, matter?

For the tool-use phase, contact-centric alternatives are also mandatory: compare system-level interface edges with a continuous contact-field condition and a contact-anchored policy frame [11, 36]. These are phase-specific baselines rather than reasons to give the graph model extra geometry or labels.

A useful diagnostic is the fraction of oracle improvement retained by inferred structure. Let M be a higher-is-better metric and M_U , M_I , and M_O the unstructured, inferred, and oracle scores. Then

$$R_{\text{oracle}} = \frac{M_I - M_U}{M_O - M_U + \epsilon}. \quad (3)$$

This ratio is descriptive, not a universal target. It is unstable when the oracle gain is small and should always be reported with raw scores and uncertainty. For error metrics, the signs must be adjusted. Its value is that it distinguishes a weak edge estimator from an unhelpful structural hypothesis.

8.3 Metrics aligned to the claimed mechanism

Contact and phase. Use per-class precision, recall, and macro-F1 for giver-only, dual-contact, receiver-only, no-contact, tool-contact, and slip states. Frame accuracy should be supplemented with transition timing error, segment-level F1, and early-warning horizon. Long stable segments otherwise dominate the score.

Edges. Dynamic-edge evaluation requires precision–recall AUC under class imbalance, onset/offset timing, false persistent edges, and calibration. Node/edge identity must be permutation-consistent where appropriate. Testing only edges used to construct training labels risks circularity; held-out contact configurations and sensor ablations are stronger.

Load and wrench. Report load-share MAE, absolute force/wrench MAE in a stated frame, angular error for force direction, normalized error across object masses, and horizon-specific future prediction. Estimated tactile wrench should not be described as independent ground truth. Gravity compensation, filtering, and coordinate transformations belong in the method, not supplementary guesswork.

Slip and uncertainty. Slip is an event-detection problem with severe imbalance. Precision–recall, false alarms per minute, and warning time are more informative than accuracy. Fine-grained contact-state and slip-direction sensing provides a concrete local classification reference [37], but local class accuracy should not be reported as evidence of cross-interface load inference. For uncertainty, use reliability diagrams, expected calibration error, Brier score, negative log likelihood, and risk–coverage curves. Calibration should be re-evaluated under object, sensor, and topology shift.

Duration, relaxation, and damage. Sustained-contact evaluation should stratify error by contact duration, temperature, material, and load/unload history; relaxation-aware force sensing shows why a single aggregate MAE can hide systematic drift [56]. For deformable objects, report task success and a separate deformation or damage constraint. SoftVTBench’s deformation-aware success rate is one example of requiring both completion and bounded peak deformation [25]; the exact threshold and evaluator-only state must remain explicit.

Closed-loop outcomes. Success should be paired with time, peak force, impulse, drop rate, object damage proxy, path error, number of recovery events, and unnecessary grip. A structured model that achieves the same success with lower peak force or earlier safe release can still be valuable. Failure-aware execution further requires separate accounting for detection, local retry, reset/recovery, and irrecoverable failure. FLARE operationalizes these branches with an online monitor and distinct retry/reset mechanisms [72]; it is a visual–language recovery framework rather than a tactile representation, but it makes recovery-aware metrics a necessary control for long-horizon claims.

8.4 Splits and out-of-distribution tests

Random frame splits are invalid for temporally correlated episodes. The split hierarchy should include, where relevant:

- episode and contiguous contact-sequence holds-out;
- object identity, geometry, mass, friction, and compliance holds-out;
- grasp pose and hand-role swap;
- collection session, operator/demonstrator, and calibration session;
- sensor instance, sensor type, missing channels, and added latency;
- embodiment or morphology;

- task/contact topology, such as handover training and tool-use testing.

The split must be made before window extraction so adjacent windows cannot cross train/test boundaries. Hyperparameter selection should use a validation split with the same unit as the intended generalization claim.

The 30-paper full-text core contained zero studies that met a strict contact-topology OOD criterion under deployment-observable inference. Strict topology OOD holds out the contact-participant or edge structure itself (for example, hand–object–hand training followed by hand–tool–workpiece–hand testing). Unseen object instances, new tool geometry, contact-mode variation within a fixed participant structure, and semantic tool–action composition are valuable tests, but they must be reported separately rather than relabelled as topology transfer.

8.5 Causal and counterfactual interventions

Correlational probes show what can be decoded; interventions show whether the claimed structure matters. Low-cost interventions include time-shuffling one hand’s tactile stream, swapping left/right roles, masking one fingertip, delaying receiver contact, perturbing the workpiece, rewiring predicted edges, or freezing edge probabilities. Physical interventions include known mass/friction changes, brief stabilizer motion, tool stiffness changes, or an unexpected contact. The expected direction of change should be stated before the experiment.

For temporal baselines, mask history content, shuffle its order, and disable observation-driven latent correction exactly rather than replacing the model with an unrelated architecture. TemporalFlow-VLA and PredVLA make these interventions especially concrete [49, 67]. For hierarchical tactile prediction, compare flat heads, stop-gradient removal, and dependency-order shuffling [64].

For example, delaying receiver contact should delay load-share crossover and safe release. If a model’s phase prediction is unchanged, it may be following a motion schedule rather than tactile evidence. Rewiring a tool–workpiece edge should damage future-wrench prediction more than rewiring an inactive relation. If all rewirings have equal effect, the graph may act as generic regularization.

8.6 Statistics and reproducibility

Formal results should use at least five independently seeded runs where training variance is material; early internal decisions can use three seeds but must not be reported as definitive. Pair results by the same data split and environment seed. Report effect sizes and confidence intervals, not only significance tests. For episode success, use bootstrap or binomial intervals with episode as the unit; for dense time series, aggregate within episode before testing or use hierarchical resampling.

Reproducibility requires a data lineage from raw sensor packet to metric: hashes or version identifiers, synchronization configuration, calibration, split manifests, normalization fitted only on training data, model configuration, seeds, software environment, checkpoint selection, and evaluation script. Code audits should distinguish a paper’s claims from what a public snapshot executes, as illustrated by the PhysGraph case in [section 6](#).

Table 8: Claim-to-evidence checklist. A paper need not include every row, but the evidence should match its stated contribution.

Claim	Minimum comparison	Stronger evidence
Touch is necessary	proprio/vision-only and tactile ablation	sensor masking and latency interventions
Explicit structure helps	matched attention/temporal/static graph	random-edge control and oracle–inferred analysis
Edges are observation-grounded	no test-time privileged inputs	pose/sensor corruption and edge calibration
Representation is physical	interface/load predictive targets	force-balance residuals and counterfactual response
Generalizes	object/session/contact-sequence split	sensor, embodiment, role, and topology holds-out
Improves control	multi-seed success and safety metrics	disturbance recovery and risk–coverage analysis

9 Synthesis and Research Agenda

9.1 What the literature currently supports

The mapped evidence supports five relatively strong conclusions.

First, tactile-specific pretraining is useful. Across robotic play, self-supervision, multimodal alignment, and cross-sensor learning, tactile representations outperform purely task-specific training in several downstream settings [12, 15, 19, 66]. This justifies reusing a strong local encoder rather than learning raw tactile features from scratch for every bimanual experiment.

Second, spatial and relational inductive biases can be useful when the task has known structure. Morphology graphs, tactile graphs, object-centric dynamics, and contact fields have succeeded in in-hand manipulation, physics prediction, packing, and contact tracking [1, 13, 18, 68]. This establishes plausibility, not universal superiority over attention.

Third, bimanual dexterity has become an active, increasingly crowded research area. Large simulation suites, demonstration-generation systems, retargeting, and real policies now cover many complex tasks [5, 24, 29, 50]. Novelty must therefore be located in a specific mechanism or evaluation gap.

Fourth, contact-centric and tactile-reactive control is moving toward broader embodiments, tools, and data scales [36, 38, 42, 70, 74]. Because many examples are recent preprints or programme-listed works, they define a competitive frontier more than a settled empirical consensus.

Fifth, the strongest competing explanations are increasingly structured even when they are not tactile graphs. Directed tactile prediction [64], mesh-prior plus diffusion-residual dynamics [57], temporal-flow supervision [67], prediction-error recurrence [49], and failure-aware recovery [72] each isolate a different source of performance. This makes matched objective, history, and recovery controls part of the minimum evidence for an explicit contact-structure claim.

9.2 What the literature does not yet establish

The evidence does not yet establish that an inferred dynamic contact graph is better than a matched temporal Transformer for bimanual load transfer. It does not establish that local tactile foundation models preserve cross-interface wrench information. It does not show that simulator contact topology can be recovered reliably from deployable touch and proprioception. It also does not provide a standard benchmark for transfer from handover to cooperative tool use.

These gaps are linked by an evidence chain:

$$\begin{aligned} \text{sensor signal} &\rightarrow \text{local contact factor} \rightarrow \text{entity/interface identity} \\ &\rightarrow \text{dynamic relation} \rightarrow \text{future physical outcome} \rightarrow \text{control decision.} \end{aligned} \quad (4)$$

Current papers often validate one or two adjacent links. A local encoder validates the first link; a privileged graph policy validates the last links conditional on correct structure; a successful bimanual imitation policy validates behaviour but may leave the intermediate mechanism unidentified. A high-impact study need not solve the complete chain, but it should state exactly which link it tests.

9.3 Eight priority problems

Table 9: Prioritized research agenda for structured bimanual tactile representations.

P	Problem	Decisive experiment	Potential deliverable
1	Observation-grounded dynamic interfaces	infer hand–object edges from tactile, proprioception, and optional vision; compare with oracle and random edges	contact-inference benchmark and calibrated edge model
2	Cross-interface physical targets	predict phase, load share, future wrench, slip, and relaxation state across short horizons and sustained contact	probe suite and physically annotated handover data
3	Matched structural comparison	compare concat, attention, directed/flat prediction, prediction-error recurrence, graph-only, residual-only, inferred, and oracle structure with equal budgets	defensible evidence for or against explicit structure
4	Cross-topology transfer	handover pretraining followed by tool–workpiece adaptation under label/data budgets	compositional transfer method and benchmark

P	Problem	Decisive experiment	Potential deliverable
5	Uncertainty under sensor failure	missing taxels, saturation, wear, thermal/viscoelastic drift, latency, and sensor-type shift	risk-aware contact representation
6	Privileged-to-observation distillation	train with simulator contacts/wrenches; deploy an observation-only student	scalable supervision without deployment leakage
7	Causal structure tests	role swap, receiver delay, history-order/edge rewiring, physical perturbation, and retry/reset separation	evidence that relations mediate predictions/control and recovery
8	Interoperable data lineage	common schema, split manifests, timestamps, frames, masks, release versions, and label provenance	reproducible BiLoad-style evidence matrix and data tooling

9.4 A staged experimental programme

The literature suggests a risk-reducing order of work.

Stage 1: audit and offline probing. If a downloadable, licensed T-Rex subset is verified, audit 20–100 episodes for timestamps, missing channels, task phases, object identity, and label provenance [42]. Otherwise, make the VTDexManip bimanual handover task the first executable benchmark while T-Rex ingestion remains a parallel artifact audit [33]. Reproduce PhysGraph as stock public code and a separately named corrected-bias variant; do not merge the two interpretations [30]. The output should be a unified schema and an initial proprio-only/concat/temporal probing table, not a large control policy.

Stage 2: representation isolation. Fix one local tactile encoder and compare matched fusion/structure variants on phase, future edge, load share, future interface wrench, slip risk, and uncertainty. The matrix should include a flat versus directed tactile hierarchy, prediction-error recurrence, training-only temporal grounding, graph-only versus global-residual-only models, and their structured combination [49, 57, 64, 67]. Use episode-, object-, and contact-sequence splits, plus contact-duration and temperature strata where sustained loading is present. Early project decisions may use three seeds, but final claims should use five or more where variance warrants. Internal go/no-go thresholds may guide resource allocation, yet they must not be written as achieved results or retrospectively optimized paper hypotheses.

Stage 3: closed-loop handover. Only after the offline model beats strong temporal baselines should it enter closed-loop quasi-static handover. Compare fixed schedules, unstructured tactile feedback, anchored tokens, and inferred structure. Evaluate safe release, peak force, drop rate, time, uncertainty, failure detection, local retry, and reset under receiver delay and object-property changes. Recovery branches should be reported separately so that controller recovery does not masquerade as a representation improvement [72].

Stage 4: gated cooperative tool use. Freeze a rigid stylus or wiping-puck–plate task early, but move to full closed-loop experiments only if (i) representation gains survive matched budgets and OOD splits, and (ii) the coordination-necessity check confirms that the stabilizing hand materially affects the outcome. The second study should emphasize transfer from a three-node handover path to a longer tool–workpiece path. If the gate fails, a valuable pivot is privileged-to-observation graph distillation or an interface-inference benchmark rather than an underdiagnosed hardware demo.

9.5 Where a bimanual contact-graph encoder can contribute

An encoder built around bimanual contact flow remains scientifically meaningful, but only under a narrow and testable formulation. Its strong contribution would not be “a GNN for two tactile hands.” It would be the combination of:

- hierarchical sensor/link/entity nodes that preserve quality masks and coordinate frames;
- dynamic, probabilistic edges inferred from deployable observation;
- object- or task-frame messages that represent load propagation through objects and tools rather than a permanent left–right shortcut;
- self-supervised or privileged-distilled edge/physics targets;

- calibrated uncertainty and explicit missing-sensor behaviour; and
- matched-budget evidence that the factorization improves transitions, data efficiency, or cross-topology generalization.

Under this formulation, the encoder is a major hypothesis and potential contribution. Without observation-grounded edges and decisive comparisons, it is an implementation variant in a crowded graph-and-touch landscape.

9.6 Publication opportunities

The research agenda naturally separates into two publishable units. The first is a benchmark-and-representation paper: synchronized handover data and probes, matched structural baselines, dynamic interface inference, and careful observability audits. The second is a transfer paper: reuse of interface/relation functions from handover in cooperative tool use, with role swap, new nodes, limited-data adaptation, and causal perturbations. A third, lower-hardware-risk alternative is an artifact paper on privileged-to-observation contact inference and reproducible code auditing across simulators.

The literature review itself could support publication after a more systematic update: preregistered database searches, deduplication across indexes, dual screening, complete author/publication metadata, code/data-availability coding, and a public evidence table. The current manuscript supplies the conceptual taxonomy and reproducible local scaffold, but it should not be labelled a systematic review until those steps are completed.

10 Limitations and Conclusion

10.1 Limitations of this review

This review is intentionally broad but not exhaustive. The archive was assembled for a concrete research decision, which biases it toward dexterous hands, learning-based control, recent tactile foundation models, and graph/contact representations. It under-samples classical grasp mechanics, human haptics, teleoperation, non-learning force control, and some industrial dual-arm systems. Screening and coding were performed by one reviewer. Five records do not have a locally validated PDF, and publication/code/data status can change rapidly. Several 2026 methods are preprints or accepted/programme-listed works; their final claims and artifacts may differ. Thirty purposively selected core papers received full-text precision coding, while the remaining 60 scholarly records retain the lighter archive-wide, rule-assisted coding layer. The precision subset is not a random sample, and neither layer used dual independent extraction or inter-rater agreement. Transparent evidence cards reduce ambiguity but do not replace those controls or independent reproduction.

The corpus is too heterogeneous for meta-analysis. Reported success rates mix different embodiments, sensors, training budgets, objects, horizons, and definitions. Counts therefore describe only the curated archive. The review also does not claim that load paths are the only useful abstraction. Deformation, compliance, friction cones, contact fields, hybrid modes, and action-conditioned latent dynamics may be superior in particular regimes.

The archive was initially frozen at 74 records on 12 August and updated to 91 on 30 August 2026 through two rapid-surveillance passes. This improves recency but does not convert the process into a database-complete systematic search. The 17 new records were integrated at the archive-wide layer; they were not retroactively treated as members of the 30-paper full-text precision subset.

10.2 Conclusion

Tactile robotics and bimanual manipulation have each advanced rapidly, but their intersection exposes a missing representational layer. Local tactile encoders identify what happens at a sensor; bimanual coordination requires a model of how several interfaces become coupled through shared entities and how that coupling changes at contact transitions. Graphs, spatial anchors, contact fields, modes, and unstructured temporal attention are competing ways to express this layer. None should be presumed best.

The updated literature makes the comparison more demanding without closing the central gap. Directed tactile prediction, sequential visuo-tactile fusion, physical-property latents, mesh-prior residual dynamics, visual world models, prediction-error recurrence, and failure-aware recovery now provide stronger alternative explanations for performance. A contact-flow model must isolate its benefit from those mechanisms and show that the benefit comes from deployable cross-interface evidence.

The strongest immediate benchmark is quasi-static handover because it provides observable phases and controlled load redistribution. Cooperative tool-workpiece manipulation is an established but complementary extension that adds a longer and role-asymmetric load path. The narrower question of matched, observation-grounded cross-hand load flow is comparatively sparse in the mapped corpus; it belongs after the simpler mechanism has passed matched-budget and coordination-necessity gates. Across both tasks, the decisive methodological requirements are deployment-observable inputs, explicit label provenance, episode/object/contact-sequence splits, oracle-inferred-random structure comparisons, uncertainty, and interventions.

The full-text core reinforces this requirement: zero of 30 papers evaluated strict held-out contact-topology transfer under deployment-observable inference. This is a scoped evidence-map result, not an estimate of the entire literature, and it motivates a decisive experiment rather than a priority claim.

An explicit bimanual contact-flow encoder can therefore be a strong contribution, but the graph layer is not the contribution. The contribution is a falsifiable factorization of sensed interfaces and object-mediated load propagation, together with evidence that this factorization improves prediction, control, or transfer under conditions where simpler temporal fusion does not. That framing provides a clear starting point for experiments and a durable standard against which future structured tactile representations can be judged.

A Review Update Record

This appendix records changes to the review corpus and interpretation after the initial manuscript freeze. It prevents three dates from being conflated: the archive snapshot used by the narrative, the completion date of the 30-paper precision audit, and the most recent surveillance date. Detailed screening notes remain in the dated files under `weekly_updates/`; this appendix records only changes that were integrated into the main manuscript.

Table 10: Version history of the review corpus and synthesis.

Date	Corpus state	Review action	Effect on interpretation
12 Aug. 2026	74 records; 69 local PDFs; five source-only	Initial narrative freeze after scope, status, and archive-wide coding	Established the six-axis taxonomy and the observation-grounded cross-interface gap.
13 Aug. 2026	30-paper purposive core: 19 Level A, 11 Level B	Completed full-text precision coding and evidence cards	Established the scoped result that 0/30 tested strict held-out contact topology under deployment-observable inference.
26 Aug. 2026	Added 10 records; archive 84; 79 local PDFs	Rapid surveillance for 20–26 Aug.; archive-wide coding and artifact audit	Added directed tactile prediction, stronger sequential fusion, physical-property probes, local slip sensing, contact-event controls, bimanual robustness/geometric baselines, and deformation-aware safety.
30 Aug. 2026	Added seven records; archive 91; 86 local PDFs	Partial-week surveillance for 27–30 Aug.; archive-wide coding and artifact audit	Added structured-prior plus global-residual dynamics, relaxation-aware force sensing, cross-embodiment/online bimanual dynamics, temporal-flow and prediction-error baselines, and explicit retry/reset recovery.
30 Aug. 2026	91 records: 90 scholarly, one resource; precision core unchanged at 30	Integrated both surveillance rounds into the main narrative and refroze the manuscript corpus	Strengthened the baseline and evaluation requirements; did not change the central gap judgment or the precision denominator.

A.1 Studies integrated from the 20–26 August update

The first surveillance round added ten studies. HiTac-WAM and VT-MUSE directly strengthened predictive and sequential visuo-tactile baselines [61, 64]. ViTacPhys, fine-grained slip sensing, and a durable tactile fingertip expanded physical-property, local-event, and hardware- reliability considerations [10, 34, 37]. Explicit physical-interaction events and Co-ToGrasp clarified the distinction between demonstrated or prescribed contact structure and topology inferred from deployment observations [39, 45]. Modality masking and PartialBiGrasp strengthened non-tactile bimanual controls [7, 26], while SoftVTBench added a safety-aware deformable-object benchmark [25].

A.2 Studies integrated from the 27–30 August update

The second round added seven studies. MeshPriorDiT motivated a three-way comparison among structured prior, expressive residual, and their hierarchy [57]; relaxation-aware soft-gripper sensing motivated contact-duration and thermal/viscoelastic drift tests [56]. Robot juggling and CLAP strengthened rapid online-dynamics and cross-embodiment video-world-model controls without supplying distributed touch [28, 32]. TemporalFlow-VLA and PredVLA motivated history-order interventions and exact open-loop prediction-error ablations [49, 67]. FLARE made failure detection, local retry, and reset/recovery separate evaluation outcomes [72].

A.3 Interpretive and artifact boundaries retained

The integrated update changes what constitutes a credible baseline, but it does not upgrade every record to full-text precision evidence. In particular:

- “contact topology” in grasp synthesis can denote a requested static layout; it is not automatically a dynamic, observation-inferred task graph;
- geometry-supervised temporal history, visual world models, and visual– language recovery remain non-tactile controls unless distributed touch is an actual deployment input;
- the ViTacPhys artifact remained forthcoming at the audit date, while the public SoftVTBench release and current manuscript described different dataset scales; and
- CLAP provided a public code/checkpoint entry, but artifact availability does not by itself establish tactile or load-topology evidence.

Future updates should append a dated row, record changes in publication and artifact status, and state explicitly whether the manuscript corpus, the precision subset, or both were changed. Counts based on the 30-paper precision audit must retain that denominator until new papers receive the same schema, located-evidence card, and quality checks.

References

- [1] Bo Ai, Stephen Tian, Haochen Shi, Yixuan Wang, Cheston Tan, Yunzhu Li, and Jiajun Wu. RoboPack: Learning tactile-informed dynamics models for dense packing. In *Proceedings of Robotics: Science and Systems*, Delft, Netherlands, July 2024. doi: 10.15607/RSS.2024.XX.130. URL <https://www.roboticsproceedings.org/rss20/p130.html>.
- [2] Kelsey R. Allen, Tatiana Lopez Guevara, Yulia Rubanova, Kim Stachenfeld, Alvaro Sanchez-Gonzalez, Peter Battaglia, and Tobias Pfaff. Graph network simulators can learn discontinuous, rigid contact dynamics. In *Proceedings of the 6th Conference on Robot Learning*, volume 205 of *Proceedings of Machine Learning Research*, pages 1157–1167. PMLR, 2023. URL <https://proceedings.mlr.press/v205/allen23a.html>.
- [3] Peter W. Battaglia, Razvan Pascanu, Matthew Lai, Danilo Rezende, and Koray Kavukcuoglu. Interaction networks for learning about objects, relations and physics. In *Advances in Neural Information Processing Systems*, volume 29, 2016. URL <https://arxiv.org/abs/1612.00222>.
- [4] Jianxin Bi, Qiang Wang, Jayaram Reddy, Kelvin Lin, Soibkhon Khajikhanov, Ruihan Gao, and Harold Soh. Heterogeneous tactile transformer, 2026. URL <https://arxiv.org/abs/2606.29948>. Preprint; under review.
- [5] Yuanpei Chen, Tianhao Wu, Shengjie Wang, Xidong Feng, Jiechuang Jiang, Stephen Marcus McAleer, Yiran Geng, Hao Dong, Zongqing Lu, Song-Chun Zhu, and Yaodong Yang. Towards human-level bimanual dexterous manipulation with reinforcement learning. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, volume 35, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/217a2a387f52c30755c37b0a73430291-Abstract-Datasets_and_Benchmarks.html.
- [6] Yuanpei Chen, Chen Wang, Yaodong Yang, and C. Karen Liu. Object-centric dexterous manipulation from human motion data. In *Proceedings of the Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, 2025. URL <https://proceedings.mlr.press/v270/chen25f.html>. Conference held in 2024; proceedings published in 2025.
- [7] Dongzhou Cheng, Ziang Li, Yixiao Zhou, Haojuan Li, Jinghao Zhang, Lei Lei, Minjing Dong, Jie Gui, and Jiaqi Wang. Robust bimanual vision-language-action models via embarrassingly simple modality masking, 2026. URL <https://arxiv.org/abs/2608.22419>. Preprint.

- [8] Trinity Chung, Kashu Yamazaki, Dhruv Patel, Alexis Duburcq, Yiling Qiao, Katerina Fragkiadaki, and Aran Nayebi. Tactile genesis: Scaling tactile simulation for robot learning, 2026. URL <https://arxiv.org/abs/2606.22332>. Preprint.
- [9] Marco Costanzo, Giuseppe De Maria, and Ciro Natale. Handover control for human-robot and robot-robot collaboration. *Frontiers in Robotics and AI*, 8:672995, 2021. doi: 10.3389/frobt.2021.672995. URL <https://doi.org/10.3389/frobt.2021.672995>.
- [10] F. Richard Cottrell, Megha H. Tippur, and Edward H. Adelson. A durable vision-based tactile fingertip for robotic manipulation, 2026. URL <https://arxiv.org/abs/2608.24242>. Preprint.
- [11] Zichen Jeff Cui, Omar Rayyan, Haritheja Etukuru, Bowen Tan, Xavier Andrianarivo, Zicheng Teng, Yihang Zhou, Krish Mehta, Nicholas Wojno, Kevin Yuanbo Wu, Manan H. Anjaria, Ziyuan Wu, Manrong Mao, Guangxun Zhang, Binit Shah, Yejin Kim, Soumith Chintala, Lerrel Pinto, and Nur Muhammad Mahi Shafiullah. Contact-anchored policies: Contact conditioning creates strong robot utility models, 2026. URL <https://arxiv.org/abs/2602.09017>. Accepted and listed in the official RSS 2026 programme.
- [12] Ruoxuan Feng, Jiangyu Hu, Wenke Xia, Tianci Gao, Ao Shen, Yuhao Sun, Bin Fang, and Di Hu. AnyTouch: Learning unified static-dynamic representation across multiple visuo-tactile sensors. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/4d893f766ab60e5337659b9e71883af4-Abstract-Conference.html.
- [13] Satoshi Funabashi, Tomoki Isobe, Fei Hongyi, Atsumu Hiramoto, Alexander Schmitz, Shigeki Sugano, and Tetsuya Ogata. Multi-fingered in-hand manipulation with various object properties using graph convolutional networks and distributed tactile sensors. *IEEE Robotics and Automation Letters*, 7(2):2102–2109, 2022. doi: 10.1109/LRA.2022.3142417. URL <https://doi.org/10.1109/LRA.2022.3142417>.
- [14] Jennifer Grannen, Yilin Wu, Brandon Vu, and Dorsa Sadigh. Stabilize to act: Learning to coordinate for bimanual manipulation. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 563–576. PMLR, 2023. URL <https://proceedings.mlr.press/v229/grannen23a.html>.
- [15] Irmak Guzey, Ben Evans, Soumith Chintala, and Lerrel Pinto. Dexterity from touch: Self-supervised pre-training of tactile representations with robotic play. In *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, pages 3142–3166. PMLR, 2023. URL <https://proceedings.mlr.press/v229/guzey23a.html>.
- [16] Jingbo He, Michael Färber, and Roberto Calandra. RCT: A robot-collected touch–vision–language dataset for tactile generalization, 2026. URL <https://arxiv.org/abs/2606.31694>. Preprint.
- [17] Liang Heng, Haoran Geng, Kaifeng Zhang, Pieter Abbeel, and Jitendra Malik. ViTacFormer: Learning cross-modal representation for visuo-tactile dexterous manipulation, 2025. URL <https://arxiv.org/abs/2506.15953>. Preprint; accepted to Robotics: Science and Systems 2026.
- [18] Carolina Higuera, Siyuan Dong, Byron Boots, and Mustafa Mukadam. Neural contact fields: Tracking extrinsic contact with tactile sensing. In *IEEE International Conference on Robotics and Automation*, pages 12576–12582. IEEE, 2023. doi: 10.1109/ICRA48891.2023.10160526. URL <https://doi.org/10.1109/ICRA48891.2023.10160526>.
- [19] Carolina Higuera, Akash Sharma, Chaithanya Krishna Bodduluri, Taosha Fan, Patrick Lancaster, Mrinal Kalakrishnan, Michael Kaess, Byron Boots, Mike Lambeta, Tingfan Wu, and Mustafa Mukadam. Sparsh: Self-supervised touch representations for vision-based tactile sensing. In *Proceedings of the 8th Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, pages 885–915. PMLR, 2025. URL <https://proceedings.mlr.press/v270/higuera25a.html>.
- [20] Carolina Higuera, Akash Sharma, Taosha Fan, Chaithanya Krishna Bodduluri, Byron Boots, Michael Kaess, Mike Lambeta, Tingfan Wu, Zixi Liu, Francois Robert Hogan, and Mustafa Mukadam. Sparsh-X: Multisensory touch representation learning, 2025. URL <https://arxiv.org/abs/2506.14754>. Preprint.
- [21] Binghao Huang, Yuanpei Chen, Tianyu Wang, Yuzhe Qin, Yaodong Yang, Nikolay Atanasov, and Xiaolong Wang. Dynamic handover: Throw and catch with bimanual hands. In *Proceedings of the Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, 2023. URL <https://proceedings.mlr.press/v229/huang23d.html>.

- [22] Jialei Huang, Yang Ye, Yuanqing Gong, Xuezhou Zhu, Yang Gao, and Kaifeng Zhang. Spatially-anchored tactile awareness for robust dexterous manipulation, 2025. URL <https://arxiv.org/abs/2510.14647>. Preprint.
- [23] Yuzhe Huang, Jiaping Wu, Jiaming Jiang, Hezhe Lin, Aikebaier Aierken, Yunlong Wang, Kun Cheng, Ziyuan Jiao, and Yuanxin Zhong. HT-Bench: Benchmarking and learning dexterous full-hand tactile representations with egocentric vision, 2026. URL <https://arxiv.org/abs/2606.19161>. Preprint.
- [24] Zhenyu Jiang, Yuqi Xie, Kevin Lin, Zhenjia Xu, Weikang Wan, Ajay Mandlekar, Linxi Fan, and Yuke Zhu. DexMimicGen: Automated data generation for bimanual dexterous manipulation via imitation learning, 2024. URL <https://arxiv.org/abs/2410.24185>. Presented at IEEE ICRA 2025.
- [25] Bowen Jing, Mingxin Wang, Ruiyang Hao, Chenchen Ge, Hanwen Shen, Junjie He, Yang Cui, Yiming Hou, Weitao Zhou, Jiawei Wang, Minglei Li, Dandan Zhang, Ding Zhao, Houde Liu, Xiaofan Li, Si Liu, Ping Luo, and Haibao Yu. SoftVTBench: A deformation-aware visuo-tactile dataset and benchmark for deformable-object manipulation, 2026. URL <https://arxiv.org/abs/2608.18701>. Preprint; public release version and manuscript scale differed at the 2026-08-30 audit.
- [26] Ayush Kaura, Vignesh Vembar, Md Faizal Karim, Keshab Patra, and K. Madhava Krishna. PartialBiGrasp: Inferring hidden local geometry for bimanual grasping from partial views, 2026. URL <https://arxiv.org/abs/2608.19188>. Preprint.
- [27] Franziska Krebs and Tamim Asfour. A bimanual manipulation taxonomy. *IEEE Robotics and Automation Letters*, 7 (4):11031–11038, 2022. doi: 10.1109/LRA.2022.3196158.
- [28] Taeyoon Lee, Chunpeng Wang, Christopher G. Atkeson, Alfred A. Rizzi, and Nicolas Rojas. Rapid on-robot learning for dynamic manipulation skills: Robot juggling, 2026. URL <https://arxiv.org/abs/2608.26800>. Preprint.
- [29] Kailin Li, Puhao Li, Tengyu Liu, Yuyang Li, and Siyuan Huang. ManipTrans: Efficient dexterous bimanual manipulation transfer via residual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025. URL https://openaccess.thecvf.com/content/CVPR2025/html/Li_ManipTrans_Efficient_Dexterous_Bimanual_Manipulation_Transfer_via_Residual_Learning_CVPR_2025_paper.html.
- [30] Runfa Blark Li, David Kim, Xinshuang Liu, Keito Suzuki, Dwait Bhatt, Nikola Raicevic, Xin Lin, Ki Myung Brian Lee, Nikolay Atanasov, and Truong Nguyen. PhysGraph: Physically-grounded graph-transformer policies for bimanual dexterous hand–tool–object manipulation, 2026. URL <https://arxiv.org/abs/2603.01436>. Preprint; under review according to the authors’ project page as checked in August 2026.
- [31] Yijiong Lin, Alex Church, Max Yang, Haoran Li, John Lloyd, Dandan Zhang, and Nathan F. Lepora. Bi-Touch: Bimanual tactile manipulation with sim-to-real deep reinforcement learning. *IEEE Robotics and Automation Letters*, 2023. URL <https://arxiv.org/abs/2307.06423>.
- [32] Kechen Liu and Ola Shorinwa. CLAP: Cross-embodiment video world models are zero-shot physical simulators, 2026. URL <https://arxiv.org/abs/2608.27406>. Preprint.
- [33] Qingtao Liu, Yu Cui, Zhengnan Sun, Gaofeng Li, Jiming Chen, and Qi Ye. VTDexManip: A dataset and benchmark for visual-tactile pretraining and dexterous manipulation with reinforcement learning. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/e19b6f65791e350347bcff8a3955cb5b-Abstract-Conference.html.
- [34] Yiwen Liu, Yujun Zhu, Kui Jia, Zhao Liao, Yangwei You, and Shuaijun Wang. ViTacPhys: Physical property-aware grasping from human visual-tactile demonstrations, 2026. URL <https://arxiv.org/abs/2608.21355>. Preprint; project page listed code and data as forthcoming at the 2026-08-30 audit.
- [35] Yun Liu, Haolin Yang, Xu Si, Ling Liu, Zipeng Li, Yuxiang Zhang, Yebin Liu, and Li Yi. TACO: Benchmarking generalizable bimanual tool-action-object understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 21740–21751, June 2024. URL https://openaccess.thecvf.com/content/CVPR2024/html/Liu_TACO_Benchmarking_Generalizable_Bimanual_Tool-Action-Object_Understanding_CVPR_2024_paper.html.
- [36] Kevin Yuchen Ma, Heng Zhang, Weisi Lin, Mike Zheng Shou, and Yan Wu. Semantic contact fields for contact-rich tool manipulation, 2026. URL <https://arxiv.org/abs/2602.13833>. Accepted and listed in the official RSS 2026 programme.

- [37] Ziyang Ma, Yuhao Sun, Zichen Ai, Xiangyang Ji, and Bin Fang. Robust slip detection and material classification via spatiotemporal transformers on a uniformly-illuminated visuo-tactile sensor, 2026. URL <https://arxiv.org/abs/2608.24162>. Accepted to IROS 2026.
- [38] Rishabh Madan, Angchen Xie, Samantha Saak, Andres Blanco, Dohyeok Lee, Sarah Grace Brown, Yunting Yan, Mark Zolotas, Jose Barreiros, and Tapomayukh Bhattacharjee. TACTIC: Tactile and vision conditioned contact-centric control for whole-arm manipulation, 2026. URL <https://arxiv.org/abs/2607.09218>. Listed in the official RSS 2026 programme.
- [39] Julien Mérand, Boris Meden, Liming Chen, and Mathieu Grossard. CoToGrasp: Contact-topology-conditioned dexterous grasp synthesis via canonical workspace learning. In *Proceedings of the European Conference on Computer Vision*, 2026. URL <https://arxiv.org/abs/2608.19776>.
- [40] Suting Ni, Hanbing Zhang, Zhenyu Wei, Guo Chen, Chixuan Zhang, Ye Shi, and Jingya Wang. TactiDex: A real-world tactile-guided benchmark for human-like dexterous manipulation, 2026. URL <https://arxiv.org/abs/2607.09190>. Preprint; author project page states acceptance to ACM Multimedia 2026, but no formal proceedings record was identified at the corpus freeze.
- [41] Yu Ning, FuQiang Zhao, and Qian Liu. TouchWGNN: Spatio-temporal tactile perception for multimodal dexterous manipulation. *Frontiers in Robotics and AI*, 13:1791424, 2026. doi: 10.3389/frobt.2026.1791424. URL <https://doi.org/10.3389/frobt.2026.1791424>.
- [42] Dantong Niu, Zhuoyang Liu, Zekai Wang, Boning Shao, Zhao-Heng Yin, Anirudh Pai, Yuvan Sharma, Stefano Saravalle, Ruijie Zheng, Jing Wang, Ryan Punamiya, Mengda Xu, Yuqi Xie, Yunfan Jiang, Letian Fu, Konstantinos Kallidromitis, Matteo Gioia, Junyi Zhang, Jiaxin Ge, Haiwen Feng, Fabio Galasso, Wei Zhan, David M. Chan, Yutong Bai, Roei Herzig, Jiahui Lei, Li Fei-Fei, Ken Goldberg, Jitendra Malik, Pieter Abbeel, Yuke Zhu, Danfei Xu, Linxi Fan, and Trevor Darrell. T-Rex: Tactile-reactive dexterous manipulation, 2026. URL <https://arxiv.org/abs/2606.17055>. Preprint.
- [43] Miquel Oller, Dmitry Berenson, and Nima Fazeli. TactileVAD: Geometric aliasing-aware tactile dynamics. In *Proceedings of the Conference on Robot Learning*, volume 229 of *Proceedings of Machine Learning Research*, 2023. URL <https://proceedings.mlr.press/v229/oller23a.html>.
- [44] Miquel Oller, Dmitry Berenson, and Nima Fazeli. Tactile-driven non-prehensile manipulation via extrinsic contact mode control. In *Proceedings of Robotics: Science and Systems*, 2024. doi: 10.15607/RSS.2024.XX.135. URL <https://www.roboticsproceedings.org/rss20/p135.html>.
- [45] A. H. G. Overbeek, H. van der Kooij, and M. Vlutters. One-shot learning from demonstration of contact-rich robotic manipulation by identifying physical interactions, 2026. URL <https://arxiv.org/abs/2608.24741>. Preprint.
- [46] Samuel Pfrommer, Mathew Halm, and Michael Posa. ContactNets: Learning discontinuous contact dynamics with smooth, implicit representations. In *Proceedings of the 2020 Conference on Robot Learning*, volume 155 of *Proceedings of Machine Learning Research*, pages 2279–2291. PMLR, 2021. URL <https://proceedings.mlr.press/v155/pfrommer21a.html>.
- [47] Hussain Sajwani, Abdulla Ayyad, Yusra Alkendi, Mohamad Halwani, Yusra Abdulrahman, Abdulqader Abusafieh, and Yahya Zweiri. TactiGraph: An asynchronous graph neural network for contact angle prediction using neuromorphic vision-based tactile sensing. *Sensors*, 23(14):6451, 2023. doi: 10.3390/s23146451. URL <https://doi.org/10.3390/s23146451>.
- [48] Alvaro Sanchez-Gonzalez, Jonathan Godwin, Tobias Pfaff, Rex Ying, Jure Leskovec, and Peter W. Battaglia. Learning to simulate complex physics with graph networks. In *Proceedings of the International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 8459–8468, 2020. URL <https://proceedings.mlr.press/v119/sanchez-gonzalez20a.html>.
- [49] Hiroki Sawada and Shunichi Kasahara. PredVLA: A sub-million-parameter predictive-coding policy for robot manipulation, 2026. URL <https://arxiv.org/abs/2608.26673>. Preprint.
- [50] Kenneth Shaw, Yulong Li, Jiahui Yang, Mohan Kumar Srirama, Ray Liu, Haoyu Xiong, Russell Mendonca, and Deepak Pathak. Bimanual dexterity for complex tasks. In *Proceedings of the Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, 2025. URL <https://proceedings.mlr.press/v270/shaw25a.html>. Conference held in 2024; proceedings published in 2025.

- [51] Yuki Shirai, Devesh K. Jha, Arvind U. Raghunathan, and Dennis Hong. Tactile tool manipulation. In *IEEE International Conference on Robotics and Automation*, pages 12597–12603. IEEE, 2023. doi: 10.1109/ICRA48891.2023.10160480. URL <https://doi.org/10.1109/ICRA48891.2023.10160480>.
- [52] Brandon Trabucco, Mariano Phielipp, and Glen Berseth. AnyMorph: Learning transferable policies by inferring agent morphology. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 21677–21691. PMLR, 2022. URL <https://proceedings.mlr.press/v162/trabucco22b.html>.
- [53] Andrea C. Tricco, Erin Lillie, Wasifa Zarin, Kelly K. O’Brien, Heather Colquhoun, Danielle Levac, David Moher, Micah D. J. Peters, Tanya Horsley, Laura Weeks, Susanne Hempel, Elie A. Akl, Christine Chang, Jessie McGowan, Lesley Stewart, Lisa Hartling, Adrian Aldcroft, Michael G. Wilson, Chantelle Garritty, Simon Lewin, Christina M. Godfrey, Marilyn T. Macdonald, Etienne V. Langlois, Karla Soares-Weiser, Jo Moriarty, Tammy Clifford, Özge Tunçalp, and Sharon E. Straus. PRISMA extension for scoping reviews (PRISMA-ScR): Checklist and explanation. *Annals of Internal Medicine*, 169(7):467–473, 2018. doi: 10.7326/M18-0850.
- [54] Mark Van der Merwe, Dmitry Berenson, and Nima Fazeli. Learning dynamics of compliant tool–environment interaction for visuo-tactile contact servoing. In *Proceedings of the Conference on Robot Learning*, volume 205 of *Proceedings of Machine Learning Research*, 2023. URL <https://proceedings.mlr.press/v205/merwe23a.html>. Conference held in 2022; proceedings published in 2023.
- [55] Tingwu Wang, Renjie Liao, Jimmy Ba, and Sanja Fidler. NerveNet: Learning structured policy with graph neural networks. In *International Conference on Learning Representations*, 2018. URL <https://openreview.net/forum?id=S1sqHMZCb>.
- [56] Yanzhe Wang, Hao Wu, Ziyi Zheng, and Huixu Dong. Relaxation-aware multimodal sensing of soft gripper driven by structure-perception-learning. In *Proceedings of Robotics: Science and Systems*, 2026. URL <https://arxiv.org/abs/2608.26622>. RSS 2026.
- [57] Zihang Wang, Jianming Hu, Shang Su, Hao Huang, Mengkai Shi, Jun Gao, and Shuo Feng. MeshPriorDiT: Hierarchical modeling for action-conditioned cloth dynamics, 2026. URL <https://arxiv.org/abs/2608.26766>. Preprint.
- [58] Zixi Wang, Zeyi Liu, Nicolas Ouporov, and Shuran Song. ContactHandover: Contact-guided robot-to-human object handover. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 9916–9923. IEEE, 2024. doi: 10.1109/IROS58592.2024.10801777. URL <https://doi.org/10.1109/IROS58592.2024.10801777>.
- [59] Youngsun Wi, Jessica Yin, Elvis Xiang, Akash Sharma, Jitendra Malik, Mustafa Mukadam, Nima Fazeli, and Tess Hellebrekers. TactAlign: Human-to-robot policy transfer via tactile alignment, 2026. URL <https://arxiv.org/abs/2602.13579>. Accepted and listed in the official RSS 2026 programme.
- [60] Longyan Wu, Jieji Ren, Chenghang Jiang, Junxi Zhou, Shijia Peng, Ran Huang, Guoying Gu, Li Chen, and Hongyang Li. TAMEn: Tactile-aware manipulation engine for closed-loop data collection in contact-rich tasks, 2026. URL <https://arxiv.org/abs/2604.07335>. Preprint.
- [61] Congsheng Xu, Qiaochu Yang, Fangyuan Shi, Yifan Han, Baijun Chen, Yiming Wang, Haonan Zhao, Daolin Ma, Xiaokang Yang, and Hesheng Wang. VT-MUSE: Multimodal unified sequential visuotactile representation learning for manipulation, 2026. URL <https://arxiv.org/abs/2608.21290>. Preprint.
- [62] Zhengtong Xu, Raghava Uppuluri, Xinwei Zhang, Cael Fitch, Philip Glen Crandall, Wan Shou, Dongyi Wang, and Yu She. UniT: Unified tactile representation for robot learning, 2024. URL <https://arxiv.org/abs/2408.06481>. Preprint.
- [63] Zhengtong Xu, Yeping Wang, Ben Abbatematteo, Jom Preechayasomboon, Sonny Chan, Nicholas Colonnese, and Amirhossein H. Memar. Contact-grounded policy learning for contact-rich manipulation, 2026. URL <https://arxiv.org/abs/2603.05687>. Accepted and listed in the official RSS 2026 programme.
- [64] Chao Xue, Chaofan Zhang, Wenxuan Ma, Guocai Yao, Shaowei Cui, and Shuo Wang. HiTac-WAM: A hierarchical tactile world action model for contact-rich robot manipulation, 2026. URL <https://arxiv.org/abs/2608.19574>. Preprint.
- [65] Fengyu Yang, Chenyang Ma, Jiacheng Zhang, Jing Zhu, Wenzhen Yuan, and Andrew Owens. Touch and go: Learning from human-collected vision and touch. In *Advances in Neural Information Processing Systems: Datasets and Benchmarks Track*, 2022. URL <https://arxiv.org/abs/2211.12498>.

- [66] Fengyu Yang, Chao Feng, Ziyang Chen, Hyoungseob Park, Daniel Wang, Yiming Dou, Ziyao Zeng, Xien Chen, Rit Gangopadhyay, Andrew Owens, and Alex Wong. Binding touch to everything: Learning unified multimodal tactile representations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/html/Yang_Binding_Touch_to_Everything_Learning_Unified_Multimodal_Tactile_Representations_CVPR_2024_paper.html.
- [67] Jiarui Yang, Yehao Lu, Yuning Su, Yu Zhong, Yufeng Xie, Yazhou Zhang, Haiyu Lan, Kaixiang Lu, Peiwen Lin, Chuang Wang, Junwei Liang, and Enyu Li. TemporalFlow-VLA: Learning physically grounded execution history for long-horizon robot manipulation, 2026. URL <https://arxiv.org/abs/2608.26821>. Preprint.
- [68] Linhan Yang, Bidan Huang, Qingbiao Li, Ya-Yen Tsai, Wang Wei Lee, Chaoyang Song, and Jia Pan. TacGNN: Learning tactile-based in-hand manipulation with a blind robot using hierarchical graph neural network. *IEEE Robotics and Automation Letters*, 8(6):3605–3612, 2023. doi: 10.1109/LRA.2023.3264759. URL <https://doi.org/10.1109/LRA.2023.3264759>.
- [69] Chengxuan Ying, Tianle Cai, Shengjie Luo, Shuxin Zheng, Guolin Ke, Di He, Yanming Shen, and Tie-Yan Liu. Do transformers really perform bad for graph representation? In *Advances in Neural Information Processing Systems*, volume 34, 2021. URL <https://proceedings.neurips.cc/paper/2021/hash/f1c1592588411002af340cbaedd6fc33-Abstract.html>.
- [70] Chengbo Yuan, Zicheng Zhang, Mingjie Zhou, Wendi Chen, Yi Wang, Zhuoyang Liu, Dantong Niu, Shuo Wang, Hui Zhang, Wenkang Zhang, Yingdong Hu, Yuanqing Gong, Wanli Xing, Chuan Wen, Cewu Lu, Kaifeng Zhang, and Yang Gao. FTP-1: A generalist foundation tactile policy across tactile sensors for contact-rich manipulation, 2026. URL <https://arxiv.org/abs/2606.13102>. Preprint.
- [71] Xinyu Zhan, Lixin Yang, Yifei Zhao, Kangrui Mao, Hanlin Xu, Zenan Lin, Kailin Li, and Cewu Lu. OAKINK2: A dataset of bimanual hands–object manipulation in complex task completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/html/Zhan_OAKINK2_A_Dataset_of_Bimanual_Hands-Object_Manipulation_in_Complex_Task_CVPR_2024_paper.html.
- [72] Ganlong Zhao, Zijia Tang, Xingping Chen, Zhanghui Kuang, Ye Tian, and Guanbin Li. FLARE: A failure-aware framework for autonomous correction and recovery in visual-language robotic manipulation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22391–22401, 2026. URL https://openaccess.thecvf.com/content/CVPR2026/html/Zhao_FLARE_A_Failure-Aware_Framework_for_Autonomous_Correction_and_Recovery_in_CVPR_2026_paper.html.
- [73] Jialiang Zhao, Yuxiang Ma, Lirui Wang, and Edward H. Adelson. Transferable tactile transformers. In *Proceedings of the Conference on Robot Learning*, volume 270 of *Proceedings of Machine Learning Research*, 2025. URL <https://proceedings.mlr.press/v270/zhao25c.html>. Conference held in 2024; proceedings published in 2025.
- [74] Jianyi Zhou, Feiyang Hong, Yunhao Li, Yicheng Zhao, Yongjue Cen, Zirui Liu, Jiakang Huang, Zirui Chen, Ruiyang Zhang, Weizhuo Zhu, Xuhua Song, and Shuo Yang. TouchWorld: A generalist tactile policy for contact-rich manipulation, 2026. URL <https://arxiv.org/abs/2607.07287>. Preprint.